



SAPIENZA
UNIVERSITÀ DI ROMA

Centralized ILP Orchestrator for Task Scheduling in Satellite Networks

Faculty of Information Engineering, Computer Science and Statistics
Computer Science

Claudio Bagini

ID number 2045337

Advisor

Dr. Emiliano Casalicchio

Co-Advisor

Academic Year 2025/2026

Thesis not yet defended

Centralized ILP Orchestrator for Task Scheduling in Satellite Networks
Bachelor's Thesis. Sapienza University of Rome

© 2026 Claudio Bagini. All rights reserved

This thesis has been typeset by L^AT_EX and the Sapthesis class.

Version: September 16, 2026

Author's email: bagini.2045337@studenti.uniroma1.it

*To countless cups of coffee, endless iterations,
and the version of me that refused to drop the task.*

Abstract

The evolution of Low Earth Orbit (LEO) satellite megaconstellations has opened new perspectives for Satellite Computing, enabling the extension of edge computing capabilities directly into space (Orbital Edge Computing). However, the highly dynamic network topology, stringent energy constraints of satellite nodes, and strict deadlines of real-time applications make task orchestration and offloading critical challenges.

This thesis presents the design, implementation, and performance evaluation of a centralized orchestrator based on Integer Linear Programming (ILP). To mitigate solver computational costs and prevent starvation, the system integrates a dynamic batching mechanism controlled by Batch Size (BS) and Batch Timeout (BT) metrics, along with a coordinated formulation of timing constraints that accounts for queuing delay prior to resolution.

Using discrete-event simulations developed with the SimPy framework, the architecture was evaluated under stress across varying request arrival rates (λ), comparing *Mapping Time* and *Mapping Energy* optimization strategies. A detailed analysis revealed that system collapse under heavy loads ($\lambda > 6$ req/s) is predominantly driven by timeout deadlines (*Deadline Exceeded*) rather than physical compute resource saturation. Furthermore, sensitivity experiments demonstrated the effectiveness of deadline relaxation (ΔD) in recovering high completion rates.

Finally, the centralized orchestrator was benchmarked against state-of-the-art distributed and heuristic solutions (*OrbitAware*, *DTS-base*, *DTS-APopt*, and distributed ILP). The results quantify the fundamental trade-off between the theoretical global optimality of centralized allocation and the impact of queuing latency compared to the responsiveness of distributed models.

Acknowledgments

Acknowledgments

Contents

1	Introduction	1
1.1	Motivation and Context	1
1.2	Problem Statement	2
1.3	Objectives and Main Contributions	3
1.4	Thesis Outline	4
2	Background and State of the Art	5
2.1	Satellite Computing Paradigms	5
2.2	Task Offloading and Routing in LEO Networks	6
2.2.1	System Entities and Network Topology	6
2.2.2	Workload Profiles and Constraints	7
2.3	Centralized vs. Distributed Orchestration	7
2.3.1	Centralized Orchestration	7
2.3.2	Distributed Orchestration	8
2.4	State-of-the-Art Approaches	8
2.4.1	Distributed Heuristic Frameworks	8
2.4.2	Optimization-Based and Distributed ILP Formulations	8
2.5	Research Gap and Thesis Positioning	9
3	System Architecture, Mathematical Formulation, and Implementation	10
3.1	Regional Constellation Architecture and Offloading Model	10
3.2	Event-Driven Dynamic Batching Engine	13
3.3	Constellation Topology Mapping and Metric Normalization	15
3.4	Centralized Integer Linear Programming Formulation	18
3.4.1	Latency, Energy, and Orbital Eclipse Modeling	18
3.4.2	Multi-Objective Function and System Constraints	19
3.4.3	Solver Routine and Cascade Rejection Diagnostics	20
3.5	Task Execution and Resilient Hybrid Routing	26
3.6	Computational Complexity and Scalability Analysis	30
3.6.1	Topology Mapping Complexity	30
3.6.2	ILP Formulation and Solver Complexity	30
3.6.3	Hybrid Routing Complexity	31

4	Performance Analysis, Optimization, and Experimental Evaluation	32
4.1	Simulation Environment and Methodological Setup	33
4.1.1	Constellation Topology and Scenario Parameters	33
4.1.2	Two-Tier Optimization Methodology	33
4.2	Macro-Tuning: Exploration and Selection of Batching Policies ($BS \times BT$)	34
4.2.1	Success Rate Dynamics	34
4.2.2	Diagnostics of Failure Causes	37
4.2.3	Response Time Decomposition	39
4.2.4	Selection of Viable Batching Archetypes	41
4.3	Micro-Tuning: Multi-Objective Weight Calibration under Stress ($\gamma \times \alpha$)	41
4.3.1	Impact on Success Rate	41
4.3.2	Failure Diagnostics and Sunset Drop Elimination	43
4.3.3	Latency Inflation and Path-Bloat Dynamics	44
4.3.4	Selection of Calibrated Objective Weights	46
4.4	Consolidated Validation: Pre- vs. Post-Optimization Benchmark . .	46
4.4.1	Success Rate Recovery	46
4.4.2	Rejection Causes Evolution	47
4.4.3	Response Time Containment	48
4.5	Summary of Findings and Architectural Readiness	49
5	Comparative Benchmark	51
5.1	Experimental Setup and Baseline Algorithms	51
5.1.1	Benchmark Dimensioning and Scenario Parameters	51
5.1.2	Evaluated Algorithms and State-of-the-Art Baselines	52
5.2	Comparative Performance Evaluation across Workload Profiles . . .	54
5.2.1	Mixed Multi-Service Workload Profile ($\beta = (0.40, 0.35, 0.25)$)	54
5.2.2	Computation-Heavy Stress: CPU-Intensive Profile ($\beta = (0, 1, 0)$)	56
5.2.3	Communication-Heavy Stress: Data-Intensive Profile ($\beta = (0, 0, 1)$)	58
5.3	Trade-Off Analysis and Operational Discussion	60
5.3.1	Theoretical Global Optimality vs. Ingress Queuing Latency .	60
5.3.2	Energy Budget Sensitivity Across Computational Dimensions	60
5.3.3	Fast-Dispatch ($BT = 0.05$ s) vs. Controlled Small-Batching ($BS = 2$)	61
5.3.4	Architectural Synthesis and Deployment Guidelines	61
5.4	Summary of Chapter Findings	61
6	Conclusions and Future Work	63
6.1	Synthesis of Contributions and Research Question Resolution	63
6.1.1	Answering the Core Research Question	63
6.1.2	Algorithmic Archetypes and Workload Sensitivity	64
6.2	Critical Limitations of the Centralized Approach	65
6.2.1	Ingress Buffering and the Batching Dilemma	65
6.2.2	Single Point of Failure and Space Hazards	65
6.2.3	State Telemetry Overhead and Information Staleness	65
6.3	Future Research Directions	66

6.3.1	Two-Tier Hybrid Orchestration Architectures	66
6.3.2	Traffic-Aware Adaptive Batching via Reinforcement Learning	66
6.3.3	Multi-Agent Distributed Reinforcement Learning (MADRL) .	66
6.3.4	Hardware-in-the-Loop (HIL) Prototyping on Space-Qualified Hardware	67
6.4	Concluding Remarks	67
Bibliography		68

Chapter 1

Introduction

In recent years, space communications and satellite networks have undergone a monumental paradigm shift. The rapid miniaturization of space components, combined with drastically reduced launch costs, has catalyzed the deployment of Mega-Constellations operating in Low Earth Orbit (LEO) [1]. Traditionally, satellite architectures functioned strictly under a "bent-pipe" model, where satellites served as passive transparent relays forwarding raw telemetry or communication signals back to Ground Stations (GS) for processing. However, the unprecedented volume of data produced by modern spaceborne applications—such as high-resolution Earth Observation (EO) [2], environmental monitoring, and ubiquitous Internet of Things (IoT) sensing [3]—has rendered this traditional pipeline highly inefficient.

To overcome the severe bandwidth limits of space-to-ground downlinks and the high propagation delays associated with remote terrestrial processing, the paradigm of **Orbital Edge Computing (OEC)**—also broadly termed Satellite Computing (SC)—has emerged [4, 5]. By equipping LEO satellites with onboard processing hardware, data can be computed directly in space close to the source, transmitting only actionable insights or highly compressed results to ground users. Despite its immense potential, executing compute-intensive tasks on resource-constrained satellite nodes introduces complex orchestration challenges that require advanced decision-making frameworks [1, 6].

1.1 Motivation and Context

The transition towards Orbital Edge Computing is driven by two main technological trends: the exponential increase in space-generated data and the strict temporal requirements of next-generation applications. Applications such as real-time disaster response, autonomous maritime navigation, and defense surveillance demand sub-second latency for task execution and reporting [2]. Processing these tasks on Earth introduces unavoidable latency bottlenecks due to intermittent Ground Station contact windows, dynamic weather interference, and severe downlink congestion.

LEO megaconstellations (such as Starlink, OneWeb, and Kuiper) consist of hundreds to thousands of satellites orbiting at altitudes between 300 km and 2000 km. These satellites move at high orbital velocities (approximately 7.5 km/s relative to the Earth's surface), resulting in a continuously changing network topology. Inter-

Satellite Links (ISLs)—utilizing high-throughput Free-Space Optics (FSO) or Radio Frequency (RF) channels—enable direct communication between adjacent satellites without routing signals down to the ground [7, 8].

Consequently, a LEO constellation effectively forms a distributed, high-velocity edge computing network in orbit. However, executing applications across this orbital fabric requires intelligent **Task Offloading and Routing Orchestration** [9, 10]. Ground users or source satellites generate computational tasks that must be dynamically assigned to suitable processing nodes across the constellation. Deciding *where* and *when* to execute a given task, while routing its associated payload across multiple dynamic ISL hops, constitutes a high-dimensional optimization problem that directly governs overall network performance.

1.2 Problem Statement

Orchestrating task offloading in Orbital Edge Computing environments presents a unique set of constraints and operational challenges that set it apart from conventional terrestrial cloud/edge systems [6]:

- **Severe Energy Budget Limits:** Satellite Edge Nodes (SENs) rely on solar panels and onboard battery storage. Computationally intensive task execution and RF/optical transmissions consume substantial power. Over-utilizing specific nodes leads to rapid battery depletion, thermal overhead, or total operational failure [6].
- **Heterogeneous and Limited Hardware Resources:** Onboard processing units (e.g., embedded CPUs, specialized AI accelerators) are constrained by strict Space, Weight, and Power (SWaP) requirements, as well as radiation-hardening constraints [5].
- **Dynamic Topology and Intermittent Connectivity:** High orbital mobility leads to frequent topology reconfigurations, changing link distances, and periodic loss of line-of-sight. Task routing paths must be continuously recalculated to avoid dropped links [7].
- **Strict Temporal Constraints and Queuing Delays:** Tasks possess strict deadlines (D_r). If the sum of transmission latency, queuing delay, and computation time exceeds D_r , the task fails, resulting in wasted system energy and degraded quality of service.

To achieve theoretical optimality in task placement and energy management, mathematical optimization tools such as **Integer Linear Programming (ILP)** are commonly employed [11, 12]. An ILP formulation can search the joint solution space of task routing, CPU allocation, and energy minimization to find globally optimal schedules.

However, centralized ILP solvers suffer from inherent computational complexity (NP-hard) [11]. In dynamic satellite networks with continuously arriving workloads, solving an ILP instance for every individual task request introduces unacceptable

computational overhead and processing latency at the central orchestrator. Conversely, accumulating task requests to solve them in batches introduces a trade-off: larger batches allow the ILP solver to achieve superior global resource allocation, but prolonged waiting times in orchestrator queues increase the risk of task deadline expiration (*Deadline Exceeded* failures).

1.3 Objectives and Main Contributions

The main objective of this thesis is to design, implement, and rigorously evaluate a **Centralized ILP Orchestration Framework** optimized for task routing and resource allocation in LEO satellite networks, incorporating a **Dynamic Batching Mechanism** to balance solver computational complexity with network queuing delay.

Specifically, this work addresses the key question: *To what extent can a centralized optimization approach maintain global allocation efficiency in high-density satellite networks before queuing and solver overhead cause systemic performance degradation?*

The main contributions of this thesis are summarized as follows:

1. **Formalization of the Centralized ILP Optimization Model:** We formulate a comprehensive ILP mathematical model that jointly optimizes task routing across Inter-Satellite Links and computational resource allocation on Satellite Edge Nodes. The model supports dynamic weighting between latency optimization (*Mapping Time*) and total energy preservation (*Mapping Energy*).
2. **Design of an Integrated Dynamic Batching & Timeout Mechanism:** To mitigate the computational overhead of solving the ILP model continuously, we integrate an adaptive batching policy governed by Batch Size (*BS*) and Batch Timeout (*BT*). Furthermore, we propose a mathematical deadline compensation formula to adjust task deadlines based on orchestrator queuing duration (T_{batch}).
3. **Design and Integration of the Centralized Orchestration Layer:** We design and implement the centralized ILP orchestration module and dynamic batching logic, integrating them into an existing Python/**SimPy** discrete-event simulation framework [13, 14]. The resulting environment accurately captures request arrival dynamics (λ), inter-satellite communication links, queue management, hardware processing, and fleet energy consumption under heterogeneous workload conditions.
4. **In-Depth Performance and Bottleneck Analysis:** We perform extensive stress testing under high workload arrival rates ($\lambda > 6$ req/s) to identify system collapse thresholds, distinguishing between physical resource saturation (CPU/Network capacity) and temporal violations (*Deadline Exceeded*). We also conduct sensitivity analyses regarding deadline relaxation (ΔD).
5. **Comparative Benchmark Against State-of-the-Art Solutions:** We evaluate the centralized ILP orchestrator against established baseline heuristics (*OrbitAware*, *DTS-base*, *DTS-APopt*) and distributed ILP formulations

from the literature [6], establishing clear operational trade-offs regarding task completion rate, overall latency, and energy efficiency.

1.4 Thesis Outline

The remainder of this thesis is structured as follows:

- **Chapter 2 (Background and State of the Art):** Reviews the fundamental concepts of Satellite Computing, LEO network dynamics, and existing literature on task offloading. It provides a taxonomy of centralized versus distributed orchestration schemes and details key reference algorithms.
- **Chapter 3 (System Architecture and ILP Formulation):** Presents the formal mathematical system model, decision variables, objective functions (*Mapping Time* vs. *Mapping Energy*), constraints, and the dynamic batching/timeout mechanisms.
- **Chapter 4 (Performance Analysis and Evaluation):** Describes the SimPy discrete-event experimental setup and evaluates the proposed centralized orchestrator under varying batching configurations, workload profiles, heavy-load conditions, and deadline relaxation scenarios.
- **Chapter 5 (Comparative Benchmark):** Details the comparative analysis benchmarking our centralized orchestrator against state-of-the-art distributed heuristics and formulations, highlighting key trade-offs in latencies, energy consumption, and global optimality.
- **Chapter 6 (Conclusions and Future Work):** Summarizes the main findings of this research, outlines limitations of centralized orbital orchestration, and proposes future research directions toward hybrid adaptive architectures.

Chapter 2

Background and State of the Art

The rapid proliferation of Low Earth Orbit (LEO) satellite constellations has fundamentally redefined the landscape of satellite communications and space-borne computing [1]. Transitioning from traditional passive data relays to active processing nodes requires a comprehensive understanding of orbital dynamics, task offloading models, and orchestration strategies [4, 5]. This chapter provides the theoretical foundation and contextual literature for this thesis. It details the paradigms of Orbital Edge Computing (OEC), analyzes task offloading and routing dynamics in LEO networks, compares centralized and distributed orchestration models, reviews state-of-the-art baseline solutions, and articulates the research gap addressed by our work.

2.1 Satellite Computing Paradigms

Space-based networking architectures have historically been categorized by their orbital altitudes: Geostationary Earth Orbit (GEO), Medium Earth Orbit (MEO), and Low Earth Orbit (LEO) [1]. GEO satellites, positioned at approximately 35,786 km, offer stationary coverage over vast geographic areas but suffer from high round-trip propagation delays (exceeding 250 ms) and high free-space path loss. MEO satellites (2,000 km – 35,786 km) reduce latency slightly but remain unsuitable for ultra-low-latency processing.

In contrast, LEO satellites operate at altitudes ranging from 300 km to 2,000 km. Their proximity to Earth drastically lowers propagation latency (10 – 40 ms round-trip) and reduces power loss, making them ideal for high-throughput, real-time communications [7, 2]. However, LEO satellites travel at extreme velocities relative to the surface (~ 7.5 km/s), completing an orbit in approximately 90 – 120 minutes. This rapid mobility results in dynamic topologies where individual satellites remain in line-of-sight with a fixed ground position for only a few minutes [3].

To process data efficiently within this dynamic environment, satellite network architectures have evolved through three primary paradigms [4, 1]:

1. **Transparent Bent-Pipe Relay:** Satellites act merely as transparent repeaters. Raw data captured by Earth Observation (EO) sensors or transmitted by Ground Users (GUs) is amplified and forwarded directly to Ground Stations (GSs). Computation occurs entirely in terrestrial cloud centers. This

model suffers from severe downlink bottlenecks and complete connectivity loss when flying over remote regions (e.g., oceans or polar zones) lacking GS infrastructure [2].

2. **Terrestrial Cloud Offloading with Inter-Satellite Routing:** Satellites utilize Inter-Satellite Links (ISLs) to route raw data across orbital planes until a satellite with an active GS connection is reached. While this mitigates local GS coverage gaps, the volume of raw data still strains optical and RF link bandwidth [7, 8].
3. **Orbital Edge Computing (OEC) / Satellite Edge Computing (SEC):** Satellites are equipped with onboard processing units—termed Satellite Edge Nodes (SENs)—capable of executing compute-intensive algorithm payloads (e.g., AI inference, image compression, signal filtering) directly in orbit [5, 10]. Raw data is processed locally at the source node or offloaded across ISLs to adjacent SENs. Only refined, actionable insights or small computation results are downlinked to Earth, drastically reducing bandwidth consumption and end-to-end response times [3, 15].

2.2 Task Offloading and Routing in LEO Networks

Task offloading in Orbital Edge Computing involves receiving task execution requests generated by Ground Users (GUs) or space-borne sensors and strategically allocating their computational execution across available SENs while routing payloads through Inter-Satellite Links (ISLs) [9, 6].

2.2.1 System Entities and Network Topology

An OEC network architecture consists of three core structural entities [6, 16]:

- **Ground Users (GUs) / Remote Sensors:** Distributed terrestrial entities that issue computational task requests r_i . Each request is characterized by a data payload size D_r^{data} (in bits), required CPU compute cycles W_r , and a strict operational deadline D_r (in seconds) [6].
- **Satellite Edge Nodes (SENs):** LEO satellites equipped with CPU cores, memory modules, and onboard energy reserves (solar panels and battery banks). Each SEN s_j possesses a processing capacity f_j (cycles/second) and an operational energy budget E_j [6, 17].
- **Inter-Satellite Links (ISLs):** High-speed communication channels connecting adjacent satellites. Intra-plane ISLs connect satellites within the same orbital plane (stable distance), whereas inter-plane ISLs connect satellites in adjacent planes (dynamically varying distance and pointing angles) [7, 8]. Each ISL link (u, v) is constrained by bandwidth B_{uv} and propagation delay τ_{uv} .

2.2.2 Workload Profiles and Constraints

Workloads submitted to an OEC network are typically heterogeneous, categorized into two main profiles [6]:

- **Real-Time Workloads:** Tasks generated by time-critical applications (e.g., emergency disaster detection, defense alerts) requiring immediate processing with tight deadlines D_r [2].
- **Batch / Delay-Tolerant Workloads:** Tasks associated with routine telemetry, scientific processing, or background data aggregation with relaxed deadlines [4].

Offloading decisions must jointly satisfy physical compute constraints ($\sum W_r \leq f_j$), network bandwidth limits, energy consumption limits (computation power P_{comp} and transmission power P_{tx}), and strict end-to-end deadline bounds ($T_{total} = T_{propagation} + T_{transmission} + T_{queue} + T_{computation} \leq D_r$) [9, 6].

2.3 Centralized vs. Distributed Orchestration

Managing resource allocation across hundreds of moving SENs requires an orchestration paradigm [1]. Literature categorizes these paradigms into two primary architectures: Centralized and Distributed.

2.3.1 Centralized Orchestration

In a centralized orchestration architecture, a dedicated master entity—such as a designated lead satellite (Ground/Orbital Orchestrator) or a ground control station—collects global network state information (CPU load, battery state, link traffic) from all SENs [1, 6].

- **Advantages:** The central orchestrator maintains complete global visibility over all network resources, enabling mathematically optimal decision-making. Global optimization models (e.g., Integer Linear Programming) can eliminate load imbalances, prevent local hotspots, and minimize network-wide energy consumption [11, 12].
- **Disadvantages:** Centralized schemes incur significant control signaling overhead to aggregate state information across the constellation. Furthermore, routing all requests to a central orchestrator creates queuing delays and bottlenecks under high request arrival rates (λ) [6]. Most critically, relying on a centralized controller introduces a Single Point of Failure (SPOF): if the orchestrator experiences hardware faults, orbital visibility loss, or link outages, the entire scheduling and task-handling pipeline is compromised, severely undermining system fault tolerance [5].

2.3.2 Distributed Orchestration

In distributed orchestration, individual SENs or localized clusters make autonomous offloading and routing decisions based solely on locally observed state information or partial neighborhood telemetry [18, 6].

- **Advantages:** Distributed schemes eliminate centralized control bottlenecks, offer sub-millisecond decision latency, and demonstrate high fault tolerance against dynamic topology changes and satellite failures [6].
- **Disadvantages:** Lacking global visibility, distributed algorithms often converge to sub-optimal solutions. They are prone to localized network congestion, unbalanced CPU utilization, and premature battery depletion on popular forwarding paths [19, 9].

2.4 State-of-the-Art Approaches

To place our research in context, we review key baseline algorithms and mathematical models established in recent IEEE literature on Orbital Edge Computing and Task Offloading.

2.4.1 Distributed Heuristic Frameworks

Heuristic approaches sacrifice theoretical optimality in favor of computational speed and low decision latency [6]:

- **OrbitAware Architecture:** A dynamic heuristic framework designed specifically for highly dynamic LEO topologies. *OrbitAware* leverages periodic orbital trajectory updates to predict link availability and select shortest processing/routing paths locally, reducing control overhead at the cost of potential link queue saturation [6].
- **DTS-base (Dynamic Task Scheduling - Base):** A greedy distributed scheduling baseline that assigns incoming tasks to the nearest available SEN with sufficient idle CPU capacity. While fast, it ignores global energy budgets and multi-hop link traffic [6].
- **DTS-APopt (Dynamic Task Scheduling - Access Point Optimized):** An enhanced version of *DTS-base* that incorporates multi-criteria decision metrics, optimizing both the local access point selection and initial ISL hop bandwidth before dispatching tasks [6].

2.4.2 Optimization-Based and Distributed ILP Formulations

To improve allocation quality, recent studies have formulated mathematical optimization models using Integer Linear Programming (ILP) and mixed-integer techniques [11, 12]:

- **Distributed / Hierarchical ILP:** Frameworks that decompose a global ILP problem into smaller sub-problems solved locally across satellite clusters [6]. While more scalable than pure centralized ILP, distributed ILP approaches suffer from convergence delays and coordination overhead during inter-cluster task transfers [17].
- **Weighted Objective Functions:** Optimization models in the literature frequently balance competing goals—such as latency minimization versus energy preservation—by assigning scalar weights (W_{lat}, W_{en}) to objective terms [9, 6].

2.5 Research Gap and Thesis Positioning

While state-of-the-art distributed heuristics offer rapid processing times and distributed ILP frameworks attempt local decomposition, a significant research gap remains regarding the practical limits of **Centralized ILP Orchestration** in LEO satellite networks [6, 1].

Existing literature frequently assumes either ideal zero-latency centralized orchestrators or dismisses centralized optimization entirely due to its computational complexity. Few works rigorously examine the trade-offs when a centralized ILP orchestrator is equipped with an integrated **Dynamic Batching** and **Timeout Mechanism** (BS/BT) to manage request queues under heavy workloads ($\lambda > 6$ req/s) [6].

This thesis addresses this exact gap by:

1. Formalizing a centralized ILP model that dynamically toggles between latency minimization (*Mapping Time*) and energy conservation (*Mapping Energy*) [6].
2. Quantifying the systemic performance bounds of centralized optimization when coupled with dynamic request batching and queuing delay compensation.
3. Conducting a rigorous comparative benchmark against established state-of-the-art heuristics (*OrbitAware*, *DTS-base*, *DTS-APopt*) and distributed ILP variants under realistic discrete-event simulation conditions (*SimPy*) [13, 6].

Chapter 3

System Architecture, Mathematical Formulation, and Implementation

This chapter provides a comprehensive dissertation on the theoretical framework, mathematical optimization formulation, and concrete software architecture of the proposed Centralized Integer Linear Programming (ILP) Orchestrator. To avoid conceptual disconnection between the theoretical modeling and its practical realization, each section progressively couples the formal system dynamics with its executable implementation in Python, utilizing the `SimPy` discrete-event simulation engine [13] and the `PuLP` mathematical programming library [14] backed by the open-source `HiGHS` solver [20].

We formalize the regional Low Earth Orbit (LEO) constellation environment, derive the multi-objective optimization problem governing task assignment and multi-hop Inter-Satellite Link (ISL) routing, detail the dynamic batching and greedy recovery mechanisms, and conclude with a rigorous computational complexity analysis.

3.1 Regional Constellation Architecture and Offloading Model

We consider a regional orbital edge computing network deployed within a Low Earth Orbit (LEO) satellite constellation [1]. The orbital network is modeled as a regional Walker-Delta LEO geometry [16]. The system comprises a set of candidate Satellite Edge Nodes (SENs), denoted by $S = \{1, 2, \dots, n\}$, and a set of regional Access Points (APs) or satellite gateways AP_{global} [6]. Within the operational region, a designated *Master* functions as the centralized orchestrator, maintaining global visibility over constellation kinematics, dynamic processor queue occupancies, residual battery energy budgets, and active inter-satellite link topologies.

Ground Users (GUs) generate computational requests that arrive dynamically at the regional gateways. Each offloading request $r \in R$ is characterized by a

five-element tuple:

$$r = (d_{cpu}, D_r, s_r, t_{arr}, AP_r) \quad (3.1)$$

where d_{cpu} is the computational CPU workload demand, D_r represents the maximum allowable service processing deadline, s_r denotes the input/output image size (in MB), t_{arr} is the generation timestamp, and $AP_r \in AP_{global}$ identifies the originating entry point [6].

To achieve robust real-time execution in highly dynamic space environments, the operational lifecycle of task r is modeled as a decoupled, three-phase circuit:

1. **Uplink Control and Forwarding Phase (T_{up}):** The originating gateway AP_r injects a lightweight control payload $S_{req} = 2$ KB into the orbital plane, which is routed across $h_{up}(r, i)$ ISL hops to candidate satellite $i \in S$ [6].
2. **On-Board Computational Execution Phase (T_{cpu}):** Satellite i processes d_{cpu} units of work, incurring both the deterministic waiting time in the on-board processor queue $W_{cpu,i}$ and the physical execution delay [4, 6].
3. **Post-Execution Geographic Result Delivery (T_{down}):** Upon computation completion, the generated output payload is routed back from the executing satellite i to the originating Access Point AP_r across the dynamic ISL topology via an autonomous geographic greedy routing protocol.

Because result return is governed dynamically in real time by the autonomous greedy protocol, the centralized ILP orchestrator focuses its optimization capacity on the joint satellite selection, uplink multi-hop routing, and processing deadline compliance.

To formalize the optimization problem and network dynamics unambiguously, Table 3.1 summarizes the foundational notations, sets, and operational parameters adopted throughout the system model and ILP formulation.

Table 3.1. Summary of mathematical notations and system parameters used in the Centralized ILP Orchestrator model.

Category	Symbol	Description
Sets & Topology	R	Set of task requests accumulated in the active batch ($ R \leq BS$).
	S	Set of candidate Satellite Edge Nodes (SENs) in the regional constellation.
	AP_{global}	Set of active regional Access Points / satellite gateways.
	AP_r	Originating Access Point for computational task request r .
	$\vec{P}_{up}(AP_r, i)$	Pre-computed shortest multi-hop ISL path from AP_r to satellite i .
Task Requests	r	Task tuple: $r = (d_{cpu,r}, D_r, s_r, t_{arr}, AP_r)$.
	$d_{cpu,r}$	Computational workload processing demand.
	D_r	Maximum allowable end-to-end service completion deadline.
	s_r	Input/output dataset payload size (in MB).
	t_{arr}	Task generation timestamp at the Access Point.
	S_{req}, S_{ref}	Uplink control payload (2 KB) and reference link payload (1 MB).
Network Links	$\tau_{u,v}$	Propagation delay of the directed link (u, v) (in seconds).
	$B_{u,v}$	Available transmission bandwidth of link (u, v) (in Bytes/s or Gbps).
	P_{net}	Radio frequency / optical laser transmission power coefficient (P_{trasm}).
	$c_{link}(u, v)$	Normalized dimensionless cost of link (u, v) in Dijkstra mapping.
	w_R, w_e	Routing cost strategy weights for delay and energy ($w_R + w_e = 1$).
	$W_{net,u}$	Network queue waiting time at intermediate forwarding node u .
Satellite Nodes	$C_{sen,i}$	Processor clock frequency / compute capacity of satellite i (Hz).
	κ_i	Dynamic CPU energy switching capacitance coefficient (e_coeff).
	B_i	Usable residual battery energy budget of satellite i (in Joules).
	$W_{cpu,i}$	Estimated queuing delay in the onboard processor execution queue.
	$t_{sunset,i}$	Solar visibility horizon / orbital eclipse timestamp for satellite i .
	Ω_i	Normalized orbital eclipse penalty factor ($\Omega_i \in [0, 1]$).
Optimization & ILP	$x_{r,i}$	Binary decision variable (1 if task r is assigned to SEN i ; 0 otherwise).
	y_r	Binary rejection variable (1 if task r is dropped by the solver; 0 otherwise).
	$R_{r,i}$	Total predicted response time for assigning task r to satellite i .
	$\epsilon_{r,i}$	Fleet cumulative energy expenditure for routing and executing task r on i .
	P_{drop}	Heavy penalty scalar applied to rejected tasks ($P_{drop} = 10^6$).
	α, β, γ	Multi-objective hierarchical weighting factors (Time, Energy, Sunset).
	BS, BT	Batch Size and Batch Timeout control thresholds for dynamic batching.

3.2 Event-Driven Dynamic Batching Engine

In a high-throughput satellite environment, invoking an NP-hard ILP solver upon each individual task arrival introduces prohibitive computational overhead and CPU starvation [11, 6]. To address this, the Master Orchestrator implements an event-driven dynamic batching mechanism inside a `SimPy` discrete-event framework [13], governed by two competing thresholds:

- **Maximum Batch Size (BS):** A structural threshold (`centralized_batch_size`) that limits the maximum number of requests solved concurrently, bounding the solver variable space $|R| \times |S|$.
- **Batch Timeout (BT):** A temporal threshold (`centralized_batch_timeout`) that enforces a maximum buffering delay, ensuring requests are not stalled indefinitely.

The execution loop remains quiescent until the arrival of the first request into `task_buffer`, which immediately succeeds `item_arrived_event` and initiates the timeout timer. The batch is dispatched when either $|Q| \geq BS$ (triggering `batch_ready_event`) or when the elapsed time satisfies $\Delta t \geq BT$.

A critical design requirement in discrete-event simulation is ensuring state consistency. Because satellite positions, CPU queue lengths, and energy levels evolve continuously according to orbital ephemeris [21], the orchestrator executes a zero-delay yield (`yield self.env.timeout(0)`) immediately prior to batch processing. This yields thread execution back to the `SimPy` event loop, ensuring that all physical updates across the constellation at timestamp t are fully committed before optimization begins.

Listing 3.1 presents the software architecture of the `Orchestrator` class initialization, request enqueueing, and the background lifecycle process.

```

1  import heapq
2  import simpy
3  import globals
4  import pulp
5  import ILP_simulation
6  import utils
7  from Task import Task
8  from EdgeServer import get_pos_proximity
9
10
11 class Orchestrator:
12     def __init__(self, env, config):
13         self.env = env
14         self.config = config
15
16         # Retrieve batch parameters from configuration
17         self.batch_size = config.get("centralized_batch_size", 5)
18         self.batch_timeout = config.get("centralized_batch_timeout", 0.25)
19
20         self.task_buffer = []
21
22         # SimPy event handlers
23         self.item_arrived_event = self.env.event()

```

```

24     self.batch_ready_event = self.env.event()
25
26     # Spawn orchestrator background process
27     self.env.process(self._run_loop())
28
29     def add_task(self, task_data, task_id, visible_aps):
30         """
31         Enqueues an incoming task into the buffer and evaluates event triggers.
32         """
33         if not self.task_buffer and not self.item_arrived_event.triggered:
34             self.item_arrived_event.succeed()
35
36         if "arrival_time" not in task_data:
37             task_data["arrival_time"] = self.env.now
38
39         self.task_buffer.append({
40             "task_id": task_id,
41             "data": task_data,
42             "visible_aps": visible_aps,
43             "arrival_time": self.env.now
44         })
45
46         if len(self.task_buffer) >= self.batch_size and not self.batch_ready_event.
↳ triggered:
47             self.batch_ready_event.succeed()
48
49     def _run_loop(self):
50         """
51         Main orchestrator lifecycle loop, handling non-blocking waits within the
↳ simulation.
52         """
53         while True:
54             # 1. Wait idle for the arrival of the FIRST task
55             if not self.task_buffer:
56                 yield self.item_arrived_event
57
58             # 2. Start batch aggregation timer (Timeout)
59             timeout_event = self.env.timeout(self.batch_timeout)
60
61             # 3. Wait until either condition is met
62             yield self.batch_ready_event | timeout_event
63
64             # 4. Process accumulated batch tasks
65             if self.task_buffer:
66                 trigger_reason = "Maximum Capacity" if self.batch_ready_event.triggered
↳ else "Timeout"
67                 print(f"[{self.env.now:.3f}] Orchestrator: ILP trigger due to {
↳ trigger_reason} ({len(self.task_buffer)} tasks).")
68
69             # --- SIMPY FIX: FORCE STATE UPDATE ---
70             # Yield control to event loop to process physical satellite state updates
71             # (CPU queues, energy, positions) at current timestamp t.
72             yield self.env.timeout(0)
73             # -----
74
75             self._process_batch()
76

```

```

77     # 5. Reset events for the subsequent cycle
78     self.item_arrived_event = self.env.event()
79     self.batch_ready_event = self.env.event()
80
81 def _process_batch(self):
82     """
83     Core routine orchestrating the resolution and dispatching pipeline.
84     """
85     # A. Map constellation topology and state
86     network_state = self._map_constellation()
87
88     # B. Invoke PuLP ILP solver
89     assignments = self._solve_ilp(network_state, self.task_buffer)
90     print(f"[{self.env.now:.3f}] Orchestrator: ILP completed.")
91
92     # C. Dispatch tasks based on optimization results
93     self._dispatch_tasks(assignments)
94
95     # Clear task buffer for subsequent batch
96     self.task_buffer.clear()

```

Listing 3.1. Orchestrator Class Architecture and SimPy Event Loop.

3.3 Constellation Topology Mapping and Metric Normalization

Before formulating the optimization model, the Master Orchestrator must compute the optimal multi-hop ISL routing paths from every active ground Access Point to all candidate Satellite Edge Nodes. This is handled by the `_map_constellation` method using a Multi-Source Dijkstra shortest-path algorithm [22].

Evaluating dynamic satellite links with raw physical dimensions introduces dimensional bias (e.g., comparing milliseconds of propagation delay against microjoules of radio energy). To achieve uniform metric evaluation across heterogeneous links, each directed edge (u, v) with latency $\tau_{u,v}$ and bandwidth $B_{u,v}$ (in Bytes/s) is evaluated against a standardized reference payload $S_{ref} = 1$ MB (1024^2 bytes) over optical ISL channels [7, 8]:

$$t_{link}(u, v) = \tau_{u,v} + \frac{S_{ref}}{B_{u,v}} \quad (3.2)$$

$$e_{link}(u, v) = P_{net} \cdot \frac{S_{ref}}{B_{u,v}} \quad (3.3)$$

where P_{net} is the radio/laser transmission power (**P**trasm) [6].

The edge metrics are normalized against theoretical upper bounds (T_{max} and E_{max}) derived from minimum network bandwidth ($B_{min} = 100$ KB/s) and maximum link latency ($\tau_{max} = 100$ ms):

$$\hat{t}_{link}(u, v) = \frac{t_{link}(u, v)}{T_{max}}, \quad \hat{e}_{link}(u, v) = \frac{e_{link}(u, v)}{E_{max}} \quad (3.4)$$

The composite dimensionless edge weight $c_{link}(u, v)$ is computed as a convex combination parameterized by weights w_R (**centralized_w_r**) and w_e (**centralized_w_e**):

$$c_{link}(u, v) = w_R \cdot \hat{t}_{link}(u, v) + w_e \cdot \hat{e}_{link}(u, v) \quad (3.5)$$

Concurrently, the routine queries all candidate nodes in `globals.edge_servers` to capture residual energy B_i , processor queue occupancy $W_{cpu,i}$ (`W_cpu_simp`), network queue occupancy $W_{net,i}$ (`W_net_simp`), processor clock speed $C_{sen,i}$, CPU energy coefficient κ_i (`e_coeff`), and orbital sunset timestamp $t_{sunset,i}$.

Listing 3.2 shows the executable implementation of the constellation mapping procedure.

```

1  def _map_constellation(self):
2      """
3      Maps shortest paths from each AP to all candidate satellites (SENs) via Dijkstra.
4      Link cost is formulated as a weighted combination of Delay and Energy.
5      """
6      # Retrieve objective cost function weights from configuration
7      w_e = self.config.get("centralized_w_e", 0.5)
8      w_R = self.config.get("centralized_w_r", 0.5)
9      p_trasm = self.config.get("Ptrasm", 1.0)
10
11     # Standardize link metrics using a 1 MB reference payload (in bytes)
12     REF_DATA_BYTES = 1024 * 1024
13
14     network_state = {
15         "routing_table": {}, # Format: {ap_name: {sen_name: {'cost': X, 'time': Y, '
↪ energy': Z, 'path': [...]}}}
16         "node_states": {}    # Format: {sen_name: {'energy_residual': X, 'cpu_queue':
↪ Y, ...}}
17     }
18
19     # 1. Collect all topology nodes to facilitate network graph traversal
20     all_nodes = globals.global_access_point + globals.edge_servers
21     node_dict = {n.name: n for n in all_nodes}
22
23     # 2. Multi-Source Dijkstra: explore network topology from each AP
24     for ap in globals.global_access_point:
25         network_state["routing_table"][ap.name] = {}
26
27         # Priority queue tuple structure:
28         # (total_cost, total_time, total_energy, current_node_name, path_so_far)
29         pq = [(0.0, 0.0, 0.0, ap.name, [ap.name])]
30
31         # Track minimum discovered cost to reach each node from this AP
32         min_costs = {ap.name: 0.0}
33
34         while pq:
35             current_cost, current_time, current_energy, current_node_name, path =
↪ heapq.heappop(pq)
36
37             # Skip obsolete (sub-optimal) popped paths
38             if current_cost > min_costs.get(current_node_name, float('inf')):
39                 continue
40
41             # If current node is a satellite (SEN), store its finalized shortest path
42             current_node_obj = node_dict[current_node_name]
43             if current_node_obj in globals.edge_servers:
44                 network_state["routing_table"][ap.name][current_node_name] = {
45                     "total_cost": current_cost,
46                     "total_time": current_time,
47                     "total_energy": current_energy,

```

```

48         "path": path
49     }
50
51     # Explore physical neighbors of current node
52     for neighbor in current_node_obj.get_neighbors():
53
54         if neighbor.name not in node_dict:
55             continue
56
57         # Extract physical communication link metrics
58         latency = current_node_obj.get_latency(neighbor) or 0.0
59         bw_Mbps = current_node_obj.get_bandwidth(neighbor) or 0.0
60         bw_Bps = bw_Mbps * (1024**2)
61
62         # Interrupted link if available bandwidth is non-positive:
63         if bw_Bps <= 0:
64             continue
65
66         # Set conservative estimates for network bandwidth lower bound and
67         latency upper bound
68         MIN_BW_BPS = 100000 # Example: 100 KB/s
69         MAX_LATENCY = 0.1 # Example: 100 ms
70
71         # Compute theoretical upper bounds for reference payload
72         MAX_LINK_TIME = MAX_LATENCY + (REF_DATA_BYTES / MIN_BW_BPS)
73         MAX_LINK_ENERGY = current_node_obj.compute_routing_energy(
74             REF_DATA_BYTES, bw_Bps, p_trasm)
75
76         # Compute latency and energy required to transmit reference payload
77         trans_time = REF_DATA_BYTES / bw_Bps
78         link_time = latency + trans_time
79         link_energy = current_node_obj.compute_routing_energy(REF_DATA_BYTES,
80             bw_Bps, p_trasm)
81
82         # Min-Max Normalization
83         norm_time = link_time / MAX_LINK_TIME if MAX_LINK_TIME > 0 else 0.0
84         norm_energy = link_energy / MAX_LINK_ENERGY if MAX_LINK_ENERGY > 0
85         else 0.0
86
87         # Composite dimensionless link cost
88         link_cost = (w_R * norm_time) + (w_e * norm_energy)
89
90         new_cost = current_cost + link_cost
91         new_time = current_time + link_time
92         new_energy = current_energy + link_energy
93
94         # Update shortest path if a lower-cost route is found
95         if new_cost < min_costs.get(neighbor.name, float('inf')):
96             min_costs[neighbor.name] = new_cost
97             heapq.heappush(pq, (new_cost, new_time, new_energy, neighbor.name
98                 , path + [neighbor.name]))
99
100     # 3. Map operational health and queue metrics of candidate SENs
101     for node in globals.edge_servers:
102         try:
103             w_cpu_simpy = node.W_cpu()
104         except Exception:

```

```

100     w_cpu_simpv = getattr(node, "W_cpu", lambda: 0.0)()
101
102     try:
103         w_net_simpv = node.W_net()
104     except Exception:
105         w_net_simpv = getattr(node, "W_net", lambda: 0.0)()
106
107     # -----
108     # Compute orbital sunset time
109     raw_sunset = getattr(node, "orbitalSunset", None)
110     if raw_sunset is not None:
111         safe_sunset = float(raw_sunset)
112     else:
113         safe_sunset = 0.0
114     # -----
115
116     # Construct node state snapshot
117     network_state["node_states"][node.name] = {
118         "energy_residual": node.energy,
119         "W_cpu_simpv": w_cpu_simpv,
120         "W_net_simpv": w_net_simpv,
121         "C_sen": getattr(node, "C_sen", 1e9),
122         "e_coeff": self.config.get("energy_coefficient", 5e-26),
123         "sunset_time": safe_sunset
124     }
125
126     return network_state

```

Listing 3.2. Constellation Topology and SEN State Mapping via Multi-Source Dijkstra.

3.4 Centralized Integer Linear Programming Formulation

Let R be the set of tasks extracted in the active batch and S the set of candidate Satellite Edge Nodes. Following classical combinatorial and linear programming principles [11, 12], the optimization problem defines two families of binary decision variables:

$$x_{r,i} \in \{0, 1\}, \quad \forall r \in R, \forall i \in S \quad (1 \text{ if task } r \text{ is allocated to satellite } i; 0 \text{ otherwise}) \quad (3.6)$$

$$y_r \in \{0, 1\}, \quad \forall r \in R \quad (1 \text{ if task } r \text{ is explicitly rejected; 0 otherwise}) \quad (3.7)$$

3.4.1 Latency, Energy, and Orbital Eclipse Modeling

For each candidate allocation (r, i) , the end-to-end processing delay $R_{r,i}$ evaluated by the solver accounts for multi-hop uplink ISL transmission across path $\vec{P}_{up}(AP_r, i)$ and local computing delay [6]:

$$R_{r,i} = \underbrace{\sum_{(u,v) \in \vec{P}_{up}} \left(W_{net,u} + \frac{S_{req}}{B_{u,v}} + \tau_{u,v} \right)}_{T_{up}(r,i)} + (W_{cpu,i} + d_{cpu,r}) \quad (3.8)$$

The cumulative fleet energy expenditure $\epsilon_{r,i}$ evaluates the network forwarding energy across all intermediate ISL nodes along $\vec{P}_{up}$ alongside local CPU execution [6]:

$$\epsilon_{r,i} = \sum_{(u,v) \in \vec{P}_{up}} P_{net} \frac{S_{req}}{B_{u,v}} + \kappa_i \cdot d_{cpu,r} \cdot C_{sen,i}^3 \quad (3.9)$$

To avoid assigning tasks to satellites nearing orbital eclipse, the orchestrator defines an orbital sunset penalty $\Omega_i \in [0, 1]$:

$$\Omega_i = \begin{cases} 1.0 - \frac{t_{sunset,i}}{\max_{k \in S}(t_{sunset,k})}, & \text{if } \max_{k \in S}(t_{sunset,k}) > 0 \\ 0.0, & \text{otherwise} \end{cases} \quad (3.10)$$

Satellites currently unreachable due to orbital eclipse ($t_{sunset,i} = 0$) receive the maximum penalty $\Omega_i = 1.0$.

3.4.2 Multi-Objective Function and System Constraints

Let $P_{drop} = 10^6$ represent a massive penalty scalar ensuring that task rejection is strictly penalized unless physical constraints are violated. The aggregated fleet latency (Z_{time}), energy (Z_{energy}), and eclipse penalty (Z_{sunset}) are defined as:

$$Z_{time} = \sum_{r \in R} \sum_{i \in S} x_{r,i} R_{r,i} + \sum_{r \in R} y_r P_{drop} \quad (3.11)$$

$$Z_{energy} = \sum_{r \in R} \sum_{i \in S} x_{r,i} \epsilon_{r,i} + \sum_{r \in R} y_r P_{drop} \quad (3.12)$$

$$Z_{sunset} = \sum_{r \in R} \sum_{i \in S} x_{r,i} \Omega_i \quad (3.13)$$

Depending on the configuration parameter `centralized_primary_objective` (Time vs. Energy), the objective function Z applies hierarchical weighting parameters (α , β , and γ) [12, 9]:

$$\min Z = \begin{cases} \alpha Z_{time} + \beta Z_{energy} + \gamma Z_{sunset}, & \text{if primary strategy is Time} \\ \alpha Z_{energy} + \beta Z_{time} + \gamma Z_{sunset}, & \text{if primary strategy is Energy} \end{cases} \quad (3.14)$$

where α denotes the weight assigned to the primary performance metric, β weights the secondary objective, and γ scales the orbital sunset penalty factor.

The optimization space is constrained by four strict linear relations:

1. **Decision Uniqueness:** Every task must be assigned to exactly one satellite or dropped:

$$\sum_{i \in S} x_{r,i} + y_r = 1, \quad \forall r \in R \quad (3.15)$$

2. **Strict Service Processing Deadline:** Enforces that uplink forwarding and local execution do not exceed the deadline D_r [6]:

$$x_{r,i} \cdot R_{r,i} \leq D_r, \quad \forall r \in R, \forall i \in S \quad (3.16)$$

3. **Distributed Fleet Energy Budget:** Enforces that the total energy allocated to satellite node k (across all tasks where k is an intermediate ISL forwarder or execution node) remains within its residual battery capacity B_k [6]:

$$\sum_{r \in R} \sum_{i \in S} x_{r,i} \cdot \epsilon_{r,i,k} \leq B_k, \quad \forall k \in S \quad (3.17)$$

4. **Constellation Load Balancing (Anti-Congestion):** Prevents processor queue bottlenecks by bounding the maximum simultaneous task allocations per node [19]:

$$\sum_{r \in R} x_{r,i} \leq \left\lfloor \frac{|R|}{|S|} \right\rfloor + 5, \quad \forall i \in S \quad (3.18)$$

3.4.3 Solver Routine and Cascade Rejection Diagnostics

The model is constructed in Python using PuLP [14] and executed via the HiGHS solver engine [20]. Upon resolution, if a task is rejected ($y_r = 1$), the orchestrator executes a cascade diagnostic evaluation across all reachable SENSs to classify the primary drop reason:

- "No_Route": No viable ISL path exists from AP_r to any satellite.
- "Deadline_Violation": Predicted response time $R_{r,i}$ exceeds D_r .
- "Sunset_Violation": Response time exceeds available solar visibility horizon $t_{sunset,i}$.
- "Energy_Exhaustion": Residual battery B_k is insufficient for multi-hop transmission or computing.
- "Load_Balancing_Rejected": Satellite reached its maximum capacity threshold.
- "Global_ILP_Conflict": Task rejected to preserve global objective optimality under shared fleet constraints.

Listing 3.3 presents the complete `_solve_ilp` implementation.

```

1  def _solve_ilp(self, network_state, tasks):
2      """
3      Open-Source Hierarchical PuLP Optimization Formulation
4      """
5      # Initialize minimization problem
6      m = pulp.LpProblem("SECMotionModel_Centralized_ILP", pulp.LpMinimize)
7
8      # Configuration parameters
9      primary_obj = self.config.get("centralized_primary_objective", "time")
10
11     # Fixed uplink control payload (2 KB)
12     S_REQ_BYTES = 2048
13     p_net = self.config.get("Ptrasm", 1.0)
14
15     R_set = tasks

```

```

16 S_set = list(network_state["node_states"].keys())
17 all_nodes = globals.global_access_point + globals.edge_servers
18 node_dict = {n.name: n for n in all_nodes}
19
20 # Decision variable definitions
21 x = pulp.LpVariable.dicts("x", ((r, i) for r in range(len(R_set)) for i in S_set)
    ↪ , cat='Binary')
22
23 y = pulp.LpVariable.dicts("y", (r for r in range(len(R_set))), cat='Binary')
24
25 # --- CONSTRAINT 1: Decision Uniqueness (Assign or Drop) ---
26 for r_idx in range(len(R_set)):
27     # Exactly one satellite assignment or drop decision per task
28     m += pulp.lpSum(x[r_idx, i] for i in S_set) + y[r_idx] == 1, f"
    ↪ Assign_Or_Drop_{r_idx}"
29
30 time_costs = {}
31 energy_costs = {}
32
33 node_energy_expressions = {k: [] for k in S_set}
34
35 # 2. Construct link metrics and operational constraints
36 for r_idx, task in enumerate(R_set):
37     task_data = task["data"]
38     ap_origin = task["visible_aps"][0].name
39     d_cpu = task_data.get('d_cpu', 0.0)
40     deadline = task_data.get('deadline', 300.0)
41     s_r_bytes = task_data.get('image_size', 0.0) * 1024 * 1024
42
43     for i in S_set:
44         route_info_up = network_state["routing_table"].get(ap_origin, {}).get(i)
45
46         if not route_info_up:
47             m += x[r_idx, i] == 0, f"NoRoute_{r_idx}_{i}"
48             time_costs[r_idx, i] = 0
49             energy_costs[r_idx, i] = 0
50             continue
51
52         path_up = route_info_up["path"]
53         num_hops_up = len(path_up) - 1
54
55         t_up = 0.0
56         e_fwd_total = 0.0
57
58         # --- Uplink Latency and Energy Formulation ---
59         for h in range(num_hops_up):
60             node_curr = node_dict[path_up[h]]
61             node_next = node_dict[path_up[h+1]]
62             lat = node_curr.get_latency(node_next) or 0.0
63
64             bw_raw = node_curr.get_bandwidth(node_next) or 0.0
65             b_isl = (bw_raw / 8.0) * (1024 ** 2) if bw_raw > 0 else 0.0
66
67             if b_isl > 0:
68                 Wn = network_state["node_states"].get(node_curr.name, {}).get("
    ↪ W_net_simp", 0.0)
69                 d_net_hop = S_REQ_BYTES / b_isl

```

```

70         t_up += Wn + d_net_hop + lat
71
72         if hasattr(node_curr, 'compute_routing_energy'):
73             e_hop_fwd = node_curr.compute_routing_energy(S_REQ_BYTES,
74 ↪ b_isl, p_net)
75         else:
76             e_hop_fwd = p_net * (S_REQ_BYTES / b_isl)
77
78         e_fwd_total += e_hop_fwd
79
80         if node_curr.name in S_set:
81             node_energy_expressions[node_curr.name].append(x[r_idx, i] *
82 ↪ e_hop_fwd)
83
84         # --- Local CPU Execution (Latency and Energy) ---
85         node_obj = node_dict[i]
86         c_sen_i = network_state["node_states"][i]["C_sen"]
87         e_coeff_i = network_state["node_states"][i]["e_coeff"]
88
89         # Latency
90         Wc = network_state["node_states"][i].get("W_cpu_simp", 0.0)
91         t_cpu_local = Wc + d_cpu
92
93         # Energy
94         if hasattr(node_obj, 'compute_execution_energy'):
95             e_cpu_local = node_obj.compute_execution_energy(d_cpu, c_sen_i, e=
96 ↪ e_coeff_i)
97         else:
98             e_cpu_local = d_cpu * e_coeff_i * (c_sen_i ** 3)
99
100         node_energy_expressions[i].append(x[r_idx, i] * e_cpu_local)
101
102         # --- Metric Aggregation and Service Deadline Constraint ---
103         time_costs[r_idx, i] = t_up + t_cpu_local# + downlink_time
104         energy_costs[r_idx, i] = e_fwd_total + e_cpu_local# + downlink_energy
105
106         # Constraint 2: Strict Service Deadline Compliance
107         R_ri_for_constraint = t_up + t_cpu_local# + downlink_time
108         m += x[r_idx, i] * R_ri_for_constraint <= deadline, f"Deadline_{r_idx}_{i}
109 ↪ )"
110
111         # Constraint 3: Distributed Fleet Battery Energy Budget
112         for k in S_set:
113             budget_k = max(0.0, network_state["node_states"][k]["energy_residual"])
114             if node_energy_expressions[k]:
115                 m += pulp.lpSum(node_energy_expressions[k]) <= budget_k, f"Energy_Limit_{
116 ↪ k}"
117
118         # --- Constraint 4: Constellation Load Balancing (Anti-Congestion) ---
119         num_satellites = len(S_set)
120
121         # Compute fleet fair-share allocation threshold
122         if num_satellites > 0:
123             fair_share = len(R_set) // num_satellites
124         else:
125             fair_share = len(R_set)

```

```

122     MAX_TASKS_PER_NODE = fair_share + 5
123
124     for i in S_set:
125         m += pulp.lpSum(x[r_idx, i] for r_idx in range(len(R_set))) <=
↪ MAX_TASKS_PER_NODE, f"Max_Capacity_{i}"
126         # -----
127
128         # --- Orbital Eclipse (Sunset) Penalty Preparation ---
129         # Determine horizon upper bound to normalize penalty within [0, 1]
130         valid_sunsets = [network_state["node_states"][k].get("sunset_time", 0.0) for k in
↪ S_set]
131         max_sunset = max(valid_sunsets + [1.0])
132
133         sunset_penalties = {}
134         for i in S_set:
135             s_time = network_state["node_states"][i].get("sunset_time", 0.0)
136
137             # Assign maximum penalty (1.0) if satellite is currently in eclipse (s_time =
↪ 0.0)
138             sunset_penalties[i] = 1.0 - (s_time / max_sunset) if max_sunset > 0 else 0.0
139             # -----
140
141         # 3. Hierarchical Multi-Objective Optimization
142         PENALTY_VALUE = 1000000.0 # Heavy penalty scalar penalizing task rejection
143
144         obj_time = pulp.lpSum(x[r_idx, i] * time_costs[r_idx, i] for r_idx in range(len(
↪ R_set)) for i in S_set) + \
145             pulp.lpSum(y[r_idx] * PENALTY_VALUE for r_idx in range(len(R_set)))
146
147         obj_energy = pulp.lpSum(x[r_idx, i] * energy_costs[r_idx, i] for r_idx in range(
↪ len(R_set)) for i in S_set) + \
148             pulp.lpSum(y[r_idx] * PENALTY_VALUE for r_idx in range(len(R_set)))
149
150         # --- Orbital Sunset Penalty Objective Term ---
151         obj_sunset = pulp.lpSum(x[r_idx, i] * sunset_penalties[i] for r_idx in range(len(
↪ R_set)) for i in S_set)
152
153         WEIGHT_PRIMARY = self.config.get("centralized_primary_weight", 1.0)
154         WEIGHT_SECONDARY = self.config.get("centralized_secondary_weight", 1.0)
155         WEIGHT_SUNSET = self.config.get("centralized_sunset_weight", 1.0)
156
157         if primary_obj == "time":
158             m += WEIGHT_PRIMARY * obj_time + WEIGHT_SECONDARY * obj_energy +
↪ WEIGHT_SUNSET * obj_sunset, "Total_Objective"
159         else:
160             m += WEIGHT_PRIMARY * obj_energy + WEIGHT_SECONDARY * obj_time +
↪ WEIGHT_SUNSET * obj_sunset, "Total_Objective"
161
162         # 4. Model Resolution via HiGHS Solver (Integrated in PuLP)
163         solver = pulp.getSolver('HiGHS', msg=False)
164         m.solve(solver)
165
166         status = pulp.LpStatus[m.status]
167
168         # -----
169         # Track aggregate energy consumption and task slot allocations for accepted tasks
170         assigned_energy_per_node = {i: 0.0 for i in S_set}

```

```

171     assigned_tasks_per_node = {i: 0 for i in S_set}
172
173     if status in ('Optimal', 'Suboptimal'):
174         for r_idx in range(len(R_set)):
175             if pulp.value(y[r_idx]) is not None and pulp.value(y[r_idx]) < 0.5: #
176                 ↪ Task accepted
177                 for i in S_set:
178                     if pulp.value(x[r_idx, i]) is not None and pulp.value(x[r_idx, i]
179                     ↪ ]) > 0.5:
180                         assigned_energy_per_node[i] += energy_costs.get((r_idx, i),
181                         ↪ 0.0)
182                         assigned_tasks_per_node[i] += 1
183
184             # -----
185
186             # Decision Extraction and Cascade Rejection Diagnostics
187             assignments = []
188             if status in ('Optimal', 'Suboptimal'):
189                 for r_idx, task in enumerate(R_set):
190
191                     # --- REJECTED TASK ---
192                     if pulp.value(y[r_idx]) and pulp.value(y[r_idx]) > 0.5:
193                         ap_origin = task["visible_aps"][0].name
194                         deadline = task["data"].get('deadline', 300.0)
195                         specific_reason = "Unknown"
196
197                         reachable_sens = [i for i in S_set if network_state["routing_table"].
198                         ↪ get(ap_origin, {}).get(i)]
199
200                         if not reachable_sens:
201                             specific_reason = "No_Route"
202                         else:
203                             # Aggregate rejection causes across all reachable candidate
204                             ↪ satellites
205                             reasons_for_rejection = []
206
207                             for i in reachable_sens:
208                                 R_ri = time_costs.get((r_idx, i), 0)
209                                 E_ri = energy_costs.get((r_idx, i), 0)
210                                 s_time_i = network_state["node_states"][i].get("sunset_time",
211                                 ↪ 0)
212
213                                 # Compute remaining battery energy after committed
214                                 ↪ allocations
215                                 actual_energy_left = network_state["node_states"][i]["
216                                 ↪ energy_residual"] - assigned_energy_per_node.get(i, 0)
217
218                                 # Cascade diagnostic evaluation for candidate satellite i
219                                 if R_ri == 0 or R_ri > deadline:
220                                     reasons_for_rejection.append("Deadline_Violation")
221                                 elif R_ri > s_time_i:
222                                     reasons_for_rejection.append("Sunset_Violation")
223                                 elif actual_energy_left < E_ri:
224                                     reasons_for_rejection.append("Energy_Exhaustion")
225                                 elif assigned_tasks_per_node.get(i, 0) >= MAX_TASKS_PER_NODE:
226                                     reasons_for_rejection.append("Load_Balancing_Rejected")
227                                 else:
228                                     reasons_for_rejection.append("Global_ILP_Conflict")

```

```

220         if reasons_for_rejection:
221             specific_reason = max(set(reasons_for_rejection), key=
↪ reasons_for_rejection.count)
222         else:
223             specific_reason = "Unknown"
224
225
226
227         print(f"[{self.env.now:.3f}] Orchestrator: Task {task['task_id']}
↪ rejected by ILP. Cause: {specific_reason}")
228
229         assignments.append({
230             "task_id": task["task_id"],
231             "data": task["data"],
232             "assigned_sen": None,
233             "routing_path_up": [],
234             "routing_path_down": [],
235             "status": f"Scartato_{specific_reason}"
236         })
237         continue
238
239     # --- ACCEPTED TASK ---
240     assigned_node = None
241     for i in S_set:
242         if pulp.value(x[r_idx, i]) and pulp.value(x[r_idx, i]) > 0.5:
243             assigned_node = i
244             break
245
246     if assigned_node:
247         ap_origin = task["visible_aps"][0].name
248         path_up = network_state["routing_table"][ap_origin][assigned_node][
↪ path"]
249         path_down = [assigned_node, ap_origin]
250
251         assignments.append({
252             "task_id": task["task_id"],
253             "data": task["data"],
254             "assigned_sen": assigned_node,
255             "routing_path_up": path_up,
256             "routing_path_down": path_down,
257             "status": "Assigned"
258         })
259     else:
260         print(f"[{self.env.now:.2f}] [ILP Centralized] Batch Infeasible (Optimization
↪ problem unsolvable). Status: {status}")
261         for task in R_set:
262             assignments.append({
263                 "task_id": task["task_id"], "data": task["data"],
264                 "assigned_sen": None, "routing_path_up": [], "routing_path_down": [],
265                 "status": "Infeasible"
266             })
267
268     return assignments

```

Listing 3.3. PuLP ILP Formulation, HiGHS Resolution, and Cascade Rejection Diagnostics.

Once the ILP optimizer produces the assignment decision matrix, the orchestrator invokes `_dispatch_tasks` to spawn independent `_route_and_execute_task` processes for each accepted request. Due to high orbital velocities in LEO constellations, planned ISL links may experience unexpected line-of-sight obstruction or degradation mid-flight. To guarantee fault tolerance, packet routing employs a dynamic hybrid routing protocol:

- Upon packet arrival at the destination Satellite Edge Node, the transmission duration and hop count are logged in `globals.gbl_generated_tasks_data`, and the process yields to the discrete-event CPU simulator (`ILP_simulation.TaskAssignment_ILP`). Finally, task status and timing metrics are synchronized in the global dataset.

```
def _dispatch_tasks(self, assignments):

    all_nodes = {n.name: n for n in (globals.global_access_point + globals.
    ↪ edge_servers)}

    for assignment in assignments:
        task_id = assignment["task_id"]
        data = assignment["data"]
        status = assignment["status"]
        dispatch_time = self.env.now

        if status != "Assigned":
            if "No_Route" in status or "Infeasible" in status:
                arrival = data.get("arrival_time", dispatch_time)
                print(f"[{dispatch_time:.3f}] Orchestrator: ILP failed ({status}) for
    ↪ Task {task_id}. Re-enqueuing into task buffer!")

                self.task_buffer.append({
                    "task_id": task_id,
                    "data": data,
                    "visible_aps": globals.global_access_point,
                    "arrival_time": arrival
                })

            if len(self.task_buffer) >= self.batch_size and not self.
    ↪ batch_ready_event.triggered:
                self.batch_ready_event.succeed()
```

```

26         continue
27
28         # --- Final Rejection (e.g., Energy Exhaustion determined by ILP) ---
29         ap = globals.global_access_point[0]
30         print(f"[{dispatch_time:.3f}] Orchestrator: Task {task_id} DROPPED ({
↪ status}).")
31         ap.record_rejected_task(
32             task_id, data["type"], dispatch_time, data["image_size"], f"ILP_{
↪ status}", data.get("d_cpu", 0)
33         )
34         continue
35
36         # --- Accepted Task: Spawn routing and execution process ---
37         self.env.process(self._route_and_execute_task(assignment, all_nodes,
↪ dispatch_time))
38
39     def _route_and_execute_task(self, assignment, all_nodes, dispatch_time):
40         task_id = assignment["task_id"]
41         data = assignment["data"]
42         sen_name = assignment["assigned_sen"]
43         sen_obj = all_nodes.get(sen_name)
44         path_up = assignment.get("routing_path_up", [])
45
46         ap_origin = path_up[0] if path_up else globals.global_access_point[0].name
47
48         # --- 1. TRANSIENT ROUTING TASK INITIALIZATION ---
49         routing_task = Task(
50             task_id=task_id,
51             current_node=ap_origin,
52             dest_node=sen_name,
53             routingInitTime=dispatch_time,
54             task_type=data["type"],
55             image_size=data["image_size"]
56         )
57         routing_task.weight = data["image_size"]
58         routing_task.visited.add(ap_origin)
59
60         start_node = all_nodes.get(ap_origin)
61         if start_node and routing_task not in start_node.tasks:
62             start_node.tasks.append(routing_task)
63
64         # --- 2. DYNAMIC NAVIGATION LOOP (Hybrid Routing: ILP -> GREEDY) ---
65         current_node_name = ap_origin
66         target_dest_name = sen_name
67         path_index = 0
68         is_greedy = False
69
70         while current_node_name != target_dest_name:
71             current_node_obj = all_nodes.get(current_node_name)
72
73             if not is_greedy:
74                 # =====
75                 # PHASE A: NOMINAL PLANNED ROUTING (ILP)
76                 # =====
77                 next_planned = path_up[path_index + 1] if path_index < len(path_up) -
↪ 1 else None
78                 next_planned_obj = all_nodes.get(next_planned) if next_planned else

```

```

79  ↪ None
80
81      if next_planned_obj and current_node_obj.get_bandwidth(
82  ↪ next_planned_obj) > 0:
83          # Active link: traverse planned ISL trajectory
84          yield self.env.process(
85              utils.sendTask(self.env, routing_task, current_node_obj,
86  ↪ next_planned_obj, "ILP_Centralized")
87          )
88          path_index += 1
89          routing_task.visited.add(next_planned)
90      else:
91          # BROKEN LINK: Trigger greedy recovery fallback!
92          print(f"[{self.env.now:.3f}] Orchestrator: Broken link at node {
93  ↪ current_node_name} (Task {task_id}). Activating GREEDY Fallback!")
94          is_greedy = True
95
96      if is_greedy:
97          # =====
98          # PHASE B: EMERGENCY RECOVERY NAVIGATION (GREEDY)
99          # =====
100      target_dest_obj = all_nodes.get(target_dest_name)
101      dest_pos = target_dest_obj.getPositionVector(globals.ist_in_conf)
102      ranker_neighbors = []
103      neighbors_list = getattr(current_node_obj, 'neighbors',
104  ↪ current_node_obj.get_neighbors())
105
106      for server in neighbors_list:
107          if current_node_obj.get_bandwidth(server) > 0:
108              neighbor_distance = get_pos_proximity(dest_pos, server.
109  ↪ getPositionVector(globals.ist_in_conf))
110              is_ap = getattr(server, 'is_acc_point', False)
111              ranker_neighbors.append((server, neighbor_distance, is_ap))
112
113      ranker_neighbors.sort(key=lambda x: (not x[2], x[1]))
114
115      best_server = None
116      for neighbor_tuple in ranker_neighbors:
117          # Loop prevention: avoid visited satellites
118          if neighbor_tuple[0].name not in routing_task.visited:
119              best_server = neighbor_tuple[0]
120              break
121
122      if best_server:
123          yield self.env.process(
124              utils.sendTask(self.env, routing_task, current_node_obj,
125  ↪ best_server, 'GREEDY_FALLBACK')
126          )
127          routing_task.visited.add(best_server.name)
128      else:
129          print(f"[{self.env.now:.3f}] Orchestrator: GREEDY Fallback failed
130  ↪ (dead-end) at node {current_node_name}. Task {task_id} dropped.")
131          routing_task.routingEndTime = self.env.now
132          routing_task.label = "DROPPED_GREEDY_DEADEND"
133          globals.gbl_tasks.append(routing_task)
134          return

```

```

128         # =====
129         # PHASE C: COMMON SAFETY CONTROLS
130         # =====
131         if getattr(routing_task, 'label', '') == 'TTL_EXPIRED':
132             print(f"[{self.env.now:.3f}] Orchestrator: Task {task_id} dropped due
↳ to expired TTL during routing.")
133             globals.gbl_tasks.append(routing_task)
134             return
135
136             current_node_name = routing_task.current_node
137
138         # --- 3. DESTINATION REACHED ---
139         transfer_time_up = self.env.now - data.get("arrival_time", dispatch_time)
140         final_hops = routing_task.hop
141         init_time = routing_task.routingInitTime
142         end_time = self.env.now
143         algos = routing_task.algorithms_used
144
145         print(f"[{self.env.now:.3f}] Orchestrator: Task {task_id} arrived at {
↳ sen_name} in {transfer_time_up:.3f}s (Hops: {final_hops})")
146
147         # --- 4. DATASET METRIC LOGGING ---
148         globals.initial_server_counter[sen_name] = globals.initial_server_counter.get
↳ (sen_name, 0) + 1
149         globals.gbl_task_hops[str(task_id)] = final_hops
150         globals.gbl_task_final_hops[str(task_id)] = final_hops
151         globals.gbl_generated_tasks_data.append({
152             "task_id": task_id, "arrival_time": data.get("arrival_time",
↳ dispatch_time), "type": data["type"],
153             "ram": data.get("required_ram", 0), "disk": data.get("required_disk", 0),
154             "image_size": data["image_size"], "exec_time": data["d_cpu"],
155             "transfer_time": transfer_time_up, "performed_transfers": final_hops,
156             "final_hops": final_hops, "execution_server": sen_name
157         })
158
159         # --- 5. SIMULATE CPU ON-BOARD COMPUTATION ---
160         yield self.env.process(
161             ILP_simulation.TaskAssignment_ILP(
162                 env=self.env, selected_server=sen_obj, task_id=task_id, image_size=
↳ data["image_size"],
163                 arrival_time_system=data.get("arrival_time", dispatch_time), num_hops
↳ =final_hops,
164                 transfer_time=transfer_time_up, task_type=data["type"], d_cpu=data["
↳ d_cpu"], D_r=data["deadline"],
165                 net_bw_override_Bps=None, result_sink={}, allow_retry=False,
166                 routing_already_charged=True, routing_energy_from=None,
↳ transfer_already_simulated=True
167             )
168         )
169
170         # --- 6. POST-EXECUTION METRIC SYNCHRONIZATION ---
171         final_task = next((t for t in globals.gbl_tasks if t.id == task_id), None)
172         if final_task:
173             final_task.hop = final_hops
174             final_task.routingInitTime = init_time
175             final_task.routingEndTime = end_time
176             final_task.algorithms_used = algos

```

```

177         else:
178             routing_task.routingEndTime = end_time
179             routing_task.label = "REJECTED_POST_ROUTING"
180             globals().gbl_tasks.append(routing_task)

```

Listing 3.4. Task Execution and Dynamic Hybrid Routing Procedure.

3.6 Computational Complexity and Scalability Analysis

To understand the computational trade-offs of centralized ILP orchestration compared to distributed heuristics, we formally analyze the algorithmic time and space complexity across all pipeline stages [6].

3.6.1 Topology Mapping Complexity

The constellation mapping routine executes Dijkstra’s algorithm [22] from each active Access Point $|AP|$ across the constellation graph $G = (V, E)$, where $V = AP_{global} \cup S$ and E is the set of active ISL and ground links. Using a binary min-heap priority queue (implemented via Python’s native `heapq` module), the asymptotic time complexity is:

$$\mathcal{O}(|AP| \cdot (|E| + |V| \log |V|)) \quad (3.19)$$

Because satellite node degree in regular grid LEO constellations (e.g., Walker delta) is bounded by physical ISL transceivers ($|E| \leq 4|S|$) [16, 7], the mapping stage scales linearly with constellation size $|S|$.

3.6.2 ILP Formulation and Solver Complexity

For a batch containing $|R|$ requests and $|S|$ candidate Satellite Edge Nodes:

- **Decision Variable Space:** The model instantiates $|R| \cdot |S|$ binary assignment variables $x_{r,i}$ and $|R|$ rejection variables y_r , yielding a total binary dimension of $\mathcal{O}(|R| \cdot |S|)$ [11].
- **Constraint Space:** The formulation constructs $|R|$ uniqueness constraints, $|R| \cdot |S|$ deadline constraints, $|S|$ fleet battery constraints, and $|S|$ load-balancing constraints, yielding an exact row complexity of:

$$\mathcal{O}(|R| \cdot |S| + |R| + 2|S|) \quad (3.20)$$

While solving general 0-1 Integer Linear Programs is NP-hard, with a worst-case time complexity of $\mathcal{O}(2^{|R| \cdot |S|})$ [11], the branch-and-cut algorithm in `HiGHS` solves the linear relaxation at each search node in polynomial time [20]. The dynamic batching threshold $BS \leq 20$ guarantees that the active problem space remains strictly bounded, maintaining mean solving times below 10 ms in practical regional operations.

3.6.3 Hybrid Routing Complexity

Under nominal operation, planned packet forwarding requires $\mathcal{O}(h_{up})$ pointer lookups along pre-computed paths. In the event of link failure, the greedy fallback evaluates the spatial distance of adjacent neighbors (k) in $\mathcal{O}(k \log k)$ time per hop. The visited set lookup operates in $\mathcal{O}(1)$ average time using hash tables, guaranteeing that emergency routing overhead remains negligible relative to optical ISL propagation delays.

Chapter 4

Performance Analysis, Optimization, and Experimental Evaluation

This chapter presents the experimental evaluation, systematic parametric tuning, and consolidated validation of the proposed Centralized Integer Linear Programming (ILP) Orchestrator [6]. While Chapter 3 established the formal mathematical optimization model and its discrete-event simulation architecture, deploying exact centralized scheduling in Low Earth Orbit (LEO) constellations introduces fundamental operational trade-offs [4, 1]. In particular, centralized optimization must balance the queuing latency introduced by task accumulation in ingress buffers against the physical dynamics of dynamic inter-satellite topologies and strict application deadlines [6, 9].

To address these challenges, this chapter follows a progressive two-stage tuning methodology concluded by an ablation validation:

1. **Stage 1: Macro-Tuning of Ingress Batching ($BS \times BT$):** We explore an exhaustive design space of 28 batching configurations across increasing traffic arrival rates ($\lambda \in [2, 10]$ req/s) and deadline relaxation thresholds (ΔD) using baseline default weights. This isolates the queuing latency bottleneck and identifies the viable operational batching archetypes.
2. **Stage 2: Micro-Tuning of Objective Weights ($\gamma \times \alpha$):** Focusing on the viable batching archetypes under heavy load ($\lambda = 8$ req/s and $\lambda = 10$ req/s), we perform a cross-sensitivity analysis between the primary objective weight (α) and the orbital sunset penalty (γ) to maximize feasibility while preventing topological path-bloat.
3. **Stage 3: Consolidated Validation (*Prima vs. Dopo*):** With the optimal objective weights determined, we validate the selected configurations across the full arrival rate spectrum ($\lambda \in [2, 10]$ req/s), directly comparing their performance against the pre-optimization baseline (*Prima vs. Dopo*).

4.1 Simulation Environment and Methodological Setup

The experimental campaign was carried out using an event-driven discrete-event simulation engine implemented in Python with `SimPy` [13], interfaced with the `PuLP` mathematical programming library [14] and solved via the open-source `HiGHS` branch-and-cut solver [20]. Numerical telemetry and array manipulations were processed using `NumPy` [23] and `SciPy` [24], while performance distributions and diagnostic curves were generated using `Matplotlib` [25]. The simulator models orbital kinematics [21], optical Inter-Satellite Links (ISLs) [7, 8], dynamic ground-to-satellite line-of-sight contact windows, and heterogeneous onboard Satellite Edge Nodes (SENs) [6].

4.1.1 Constellation Topology and Scenario Parameters

The orbital network is modeled as a regional Walker-Delta LEO constellation operating at an altitude of 600 km with an inclination of 53° [16]. Satellites are interconnected via full-duplex optical ISLs characterized by nominal transmission bandwidths of $B_{u,v} = 1$ Gbps and propagation delays $\tau_{u,v}$ computed dynamically from orbital positions [8]. A designated Master satellite in the operational region collects offloaded task requests forwarded by Ground Users (GUs) through regional Ground Access Points (AP_{global}) [6].

Workload demands and request dimensions are parameterized through three profiling vectors, consistent with recent benchmarking literature [6]:

- **Workload Distribution Vector** $\beta = (\beta_G, \beta_C, \beta_{CD})$: Dictates the proportion of Generic tasks, CPU-intensive tasks, and Data-intensive tasks. In the baseline batching and tuning experiments, the default multi-service profile is set to $\beta = (0.40, 0.35, 0.25)$ [6].
- **Real-Time Payload Vector** $\alpha = (\alpha_M, \alpha_L, \alpha_V)$: Governs the distribution of input/output dataset dimensions across Medium (1–5 MB), Large (5–20 MB), and Very Large (20–50 MB) files, parameterized as $\alpha = (0.30, 0.50, 0.20)$ [6, 2].
- **Batch Task Data Vector** $\gamma_{data} = (\gamma_L, \gamma_V)$: Controls delay-tolerant background transfers, fixed to $\gamma_{data} = (0.70, 0.30)$ [4].

Candidate edge nodes are assigned operational battery reserves of $E_i \in \{40, 80\}$ kJ per orbital pass [6, 17]. The request arrival rate λ varies systematically from 2 req/s (low-load, mean inter-arrival time of 0.50 s) up to 10 req/s (heavy-load stress regime, mean inter-arrival time of 0.10 s) [6].

4.1.2 Two-Tier Optimization Methodology

Directly testing all permutations of batching thresholds, objective weights, and network loads would require thousands of simulations and obscure the physical origins of performance bottlenecks. We therefore decouple the optimization process:

- **Architectural Decoupling**: Ingress batching (BS, BT) governs the static waiting time of tasks in the coordinator queue *before* the ILP model is instantiated. Conversely, objective function weights (α, γ) guide the mathematical branch-and-bound trade-offs *during* solver execution [11, 12].

- **Sequential Tuning:** By first sweeping $BS \times BT$ under default weights, we isolate the feasible batching configurations that do not cause queue-induced deadline collapse. Next, we tune the objective weights (α, γ) under heavy stress on these selected batching archetypes. Finally, we benchmark the optimized configurations against the pre-tuning baseline across the entire workload spectrum.

4.2 Macro-Tuning: Exploration and Selection of Batching Policies ($BS \times BT$)

In a centralized orchestrator, incoming requests are buffered until either the Batch Size (BS) or the Batch Timeout (BT) threshold is triggered. This introduces a fundamental trade-off: larger batches enlarge the joint solution space, enabling the solver to achieve higher allocation efficiency, but prolonged buffering erodes task deadline margins [6].

To quantify this trade-off, a parametric sweep of 28 configurations was executed across arrival rates $\lambda \in [2, 10]$ req/s under both Energy Mapping ($w_{lat} = 0, w_{en} = 1$) and Time Mapping ($w_{lat} = 1, w_{en} = 0$), evaluated across deadline relaxation levels $\Delta D \in \{10\%, 20\%, 30\%\}$.

4.2.1 Success Rate Dynamics

Figures 4.1 and 4.2 show the Task Success Rate across the parameter space for Energy and Time mapping, respectively.

4.2 Macro-Tuning: Exploration and Selection of Batching Policies ($BS \times BT$)




Figure 4.1. Success Rate across arrival rates under Energy Mapping for BS -sweep ($BT = 2.0$ s) and BT -sweep ($BS = 20$) across deadline relaxation levels $\Delta D \in \{10\%, 20\%, 30\%\}$.

4.2 Macro-Tuning: Exploration and Selection of Batching Policies ($BS \times BT$) 36




Figure 4.2. Success Rate across arrival rates under Time Mapping for BS -sweep ($BT = 2.0$ s) and BT -sweep ($BS = 20$) across deadline relaxation levels $\Delta D \in \{10\%, 20\%, 30\%\}$.

The empirical trends highlight three critical dynamics:

- **Performance Collapse in High-Timeout Regimes:** Configurations relying on long buffering timeouts (e.g., $BS = 20$ with $BT \geq 0.4$ s, orange and yellow curves) exhibit severe degradation, dropping below 55% at $\lambda = 2$ req/s and falling under 30% at $\lambda = 10$ req/s.
- **Dominance of Fast-Dispatch Policies:** Configurations that enforce immediate dispatch—such as $BS = 1$, $BS = 2$ with $BT = 2.0$ s, or $BS = 20$ with $BT \leq 0.05$ s—achieve peak completion rates exceeding 90% at low arrival rates ($\lambda \leq 3$ req/s) and maintain higher resilience under stress [6].
- **Mapping Strategy Invariance:** The success curves for Time Mapping and Energy Mapping are virtually identical (differing by less than 2.5%), confirming

4.2 Macro-Tuning: Exploration and Selection of Batching Policies ($BS \times BT$) 37

that under heavy traffic, system throughput is bounded primarily by queue dwell times rather than the objective metric.

4.2.2 Diagnostics of Failure Causes

Figures 4.3 and 4.4 illustrate the relative distribution of rejection causes across the parameter sweeps.




Figure 4.3. Rejection causes distribution across arrival rates under Energy Mapping, comparing BS -sweep ($BT = 2.0$ s) against BT -sweep ($BS = 20$) for $\Delta D \in \{10\%, 20\%, 30\%\}$.

4.2 Macro-Tuning: Exploration and Selection of Batching Policies ($BS \times BT$) 38




Figure 4.4. Rejection causes distribution across arrival rates under Time Mapping, comparing BS -sweep ($BT = 2.0$ s) against BT -sweep ($BS = 20$) for $\Delta D \in \{10\%, 20\%, 30\%\}$.

The failure diagnostics reveal a fundamental structural divide:

- In fast-dispatch configurations ($BS = 1$ or $BS = 20$ with $BT \leq 0.05$ s), task rejections stem primarily from physical capacity limits: *Insufficient CPU* (green) and *Insufficient CPU+NET* (orange). Tasks bypass ingress buffering, meaning drops occur only when onboard compute or link bandwidth is saturated.
- In delayed-batch configurations ($BT \geq 0.2$ s or large BS with high timeout), *Deadline Exceeded* (blue) rapidly dominates, accounting for 70%–85% of all rejections. The orchestrator buffers tasks until their remaining slack time $\Delta t_{rem} = D_r - t_{queue}$ is insufficient to cover propagation and execution.

4.2.3 Response Time Decomposition

Figures 4.5 and 4.6 show the end-to-end response time decomposed into System Execution Time (T_{exec}) and Network/Wait Time ($T_{net+wait}$).




Figure 4.5. Response Time decomposition under Energy Mapping across arrival rates for BS -sweep and BT -sweep configurations under $\Delta D \in \{10\%, 20\%, 30\%\}$.

4.2 Macro-Tuning: Exploration and Selection of Batching Policies ($BS \times BT$)40




Figure 4.6. Response Time decomposition under Time Mapping across arrival rates for BS -sweep and BT -sweep configurations under $\Delta D \in \{10\%, 20\%, 30\%\}$.

These plots validate a central property of the centralized orchestrator:

$$T_{exec} \approx \text{const} \leq 0.60 \text{ ms}, \quad \text{while} \quad T_{net+wait} \propto \frac{1}{\mu - \lambda} \quad (4.1)$$

The solver runtime (T_{exec}) remains virtually constant regardless of the arrival rate, demonstrating that solving instances of bounded size ($|R| \leq 20, |S| \leq 15$) is computationally practical on space-qualified hardware using efficient branch-and-cut solvers [20, 11]. The latency growth observed under high traffic is driven almost entirely by the queuing component ($T_{net+wait}$), which rises from ≈ 0.40 ms to over 1.70 ms.

4.2.4 Selection of Viable Batching Archetypes

The sweep results rule out all batching policies with $BT \geq 0.2$ s or $BS \geq 8$, as queue dwell times systematically violate task deadlines. To represent the viable operational design space, we select two complementary configurations for subsequent weight calibration:

1. **Reactive Fast-Dispatch** ($BS = 20, BT = 0.05$ s): Operates with an aggressive 50 ms timeout. Because inter-arrival times at $\lambda \leq 10$ req/s are ≥ 100 ms, the timeout triggers almost continuously, dispatching 1–2 tasks per cycle and keeping buffer dwell times near zero.
2. **Controlled Small-Batching** ($BS = 2, BT = 2.0$ s): Enforces minimal task aggregation. At arrival rates $\lambda \geq 2$ req/s, the batch size threshold ($BS = 2$) is reached within 0.2–0.5 s, providing the solver with a joint optimization horizon over 2 requests without severe queue buildup.

4.3 Micro-Tuning: Multi-Objective Weight Calibration under Stress ($\gamma \times \alpha$)

Having isolated the viable batching configurations ($BS = 2, BT = 2.0$ s and $BS = 20, BT = 0.05$ s), we calibrate the scalarization weights of the multi-objective function [12]:

$$\min Z = \begin{cases} \alpha Z_{time} + \beta Z_{energy} + \gamma Z_{sunset}, & \text{if primary objective is Time} \\ \alpha Z_{energy} + \beta Z_{time} + \gamma Z_{sunset}, & \text{if primary objective is Energy} \end{cases} \quad (4.2)$$

where $\beta = 1.0$, the primary objective weight is evaluated across $\alpha \in \{10, 50, 1000\}$, and the orbital sunset deterrent penalty is evaluated across $\gamma \in \{2, 4, 6, 8, 10, 1000, 5000, 10000\}$. The cross-sensitivity experiment was conducted under heavy stress loads of $\lambda = 8$ req/s and $\lambda = 10$ req/s, comparing both Dijkstra-Time and Dijkstra-Energy link routing metrics [22, 6].

4.3.1 Impact on Success Rate

Figures 4.7 and 4.8 depict the Task Success Rate across the $(\gamma \times \alpha)$ parameter space for $\lambda = 8$ req/s and $\lambda = 10$ req/s, respectively.

4.3 Micro-Tuning: Multi-Objective Weight Calibration under Stress ($\gamma \times \alpha$) 42




Figure 4.7. Cross-sensitivity Success Rate ($\gamma \times \alpha$) at $\lambda = 8$ req/s across batching configurations, Dijkstra metrics, and mapping objectives.




Figure 4.8. Cross-sensitivity Success Rate ($\gamma \times \alpha$) at $\lambda = 10$ req/s across batching configurations, Dijkstra metrics, and mapping objectives.

The empirical trends indicate:

- **Dominance of Primary Weight $\alpha = 1000$:** In configurations sensitive to schedule optimization (such as Time | $BS : 2, BT : 2.0$), setting $\alpha = 1000$ (green bars) yields an 80% completion rate at $\lambda = 8$ req/s and 70% at $\lambda = 10$ req/s, substantially outperforming $\alpha = 10$ and $\alpha = 50$ (50%–55%). A large primary weight forces the solver to prioritize deadline compliance and physical capacity

constraints over secondary trade-offs.

- **Optimal Sunset Window $\gamma \in [6, 10]$:** Peak performance across all configurations consistently falls within $\gamma \in [6, 10]$, reaching its maximum at $\gamma = 8$ and $\gamma = 10$.
- **Performance Degradation at $\gamma \geq 1000$:** When $\gamma \geq 1000$, the Success Rate drops to $\approx 50\%$. An excessively large sunset penalty over-penalizes the objective space, causing the solver to reject feasible tasks out of an over-conservative avoidance of nodes nearing eclipse.

4.3.2 Failure Diagnostics and Sunset Drop Elimination

Figures 4.9 and 4.10 show the failure cause distributions across the $(\gamma \times \alpha)$ test matrix.




Figure 4.9. Distribution of rejection causes across the $(\gamma \times \alpha)$ parameter space at $\lambda = 8$ req/s.




Figure 4.10. Distribution of rejection causes across the $(\gamma \times \alpha)$ parameter space at $\lambda = 10$ req/s.

The diagnostics demonstrate the role of the sunset parameter:

- At low values ($\gamma \leq 4$), tasks are occasionally routed to satellites that lose line-of-sight during execution, resulting in *Sunset* drops (brown segments).
- For $\gamma \geq 8$, drops due to orbital sunset are completely eliminated.
- However, setting $\gamma \geq 1000$ increases *Deadline Exceeded* failures, as the solver selects longer paths to avoid satellites nearing the eclipse boundary.

4.3.3 Latency Inflation and Path-Bloat Dynamics

Figures 4.11 and 4.12 present the response time decomposition across the $(\gamma \times \alpha)$ configurations.

4.3 Micro-Tuning: Multi-Objective Weight Calibration under Stress ($\gamma \times \alpha$) 45




Figure 4.11. Response Time decomposition across the ($\gamma \times \alpha$) parameter space at $\lambda = 8$ req/s.




Figure 4.12. Response Time decomposition across the ($\gamma \times \alpha$) parameter space at $\lambda = 10$ req/s.

These results illustrate the path-bloat phenomenon:

- Within $\gamma \in [2, 10]$, the end-to-end response time is stable at $\approx 0.65\text{--}0.70$ s, with solver computation taking < 0.30 s.
- For $\gamma \geq 1000$, the response time surges to $1.05\text{--}1.15$ s—a 55% increase. To avoid the large penalty term Ω_i , the solver routes packets through circuitous multi-hop ISL paths or offloads them to distant satellites, inflating network propagation delay and intermediate queuing.

4.3.4 Selection of Calibrated Objective Weights

Correlating the success rate, failure distributions, and latency metrics identifies the optimal operating parameter set:

$$\alpha^* = \mathbf{1000}, \quad \gamma^* = \mathbf{8} \quad (\text{with } \gamma = \mathbf{10} \text{ as an equivalent alternative}) \quad (4.3)$$

This combination provides a balanced operational setting: $\alpha = 1000$ ensures that physical deadlines and compute capacity dominate solver decisions, while $\gamma = 8$ prevents assignments to satellites entering eclipse without inflating path lengths.

4.4 Consolidated Validation: Pre- vs. Post-Optimization Benchmark

To validate the multi-objective calibration and quantify the performance gains achieved over the initial baseline, we execute a comparative ablation study. We re-evaluate the two selected batching configurations across the entire arrival rate spectrum ($\lambda \in [2, 10]$ req/s) under Dijkstra-Time routing with $\Delta D = 10\%$, comparing:

- **Pre-Optimization Baseline (*Prima*):** The initial configuration using default objective weights ($\alpha = 10, \gamma = 2$) as evaluated during the macro-tuning sweep.
- **Post-Optimization Model (*Dopo*):** The calibrated configuration integrating the tuned weights ($\alpha = 1000, \gamma = 8$) derived from the high-load cross-sensitivity analysis.

4.4.1 Success Rate Recovery

Figure 4.13 presents the direct comparison of Task Success Rate between the initial (*Prima*) and optimized (*Dopo*) configurations under both Energy Mapping and Time Mapping.




Figure 4.13. Success Rate comparison: Pre-optimization baseline (*Prima*) vs. Post-optimization model (*Dopo*) for $BS = 2, BT = 2.0$ s and $BS = 20, BT = 0.05$ s under Dijkstra-Time routing ($\Delta D = 10\%$).

The empirical comparison confirms substantial throughput improvements across all arrival rates:

- **Significant Gain in Batching Configurations:** For $BS = 2, BT = 2.0$ s under Time Mapping, the optimized weights increase the Success Rate from 74% to 90% at $\lambda = 2$ req/s, and from 47% to 68% under heavy load ($\lambda = 10$ req/s), providing an average gain of +21% across the entire traffic range.
- **Resilience Under Heavy Stress:** For the reactive fast-dispatch setting ($BS = 20, BT = 0.05$ s), the success rate under peak load ($\lambda = 10$ req/s) improves from 38% to 43% under Time Mapping and from 39% to 41% under Energy Mapping.
- **Verification of Weight Scaling:** The cross-sensitivity tuning conducted at $\lambda = 8$ req/s and $\lambda = 10$ req/s transfers effectively across low and moderate traffic regimes ($\lambda \leq 4$ req/s), proving that prioritizing the primary objective ($\alpha = 1000$) maintains stable scheduling without overfitting to heavy-load scenarios.

4.4.2 Rejection Causes Evolution

Figure 4.14 illustrates the evolution of failure cause distributions between the pre- and post-optimization models.




Figure 4.14. Rejection causes distribution comparison: Pre-optimization baseline (*Prima*) vs. Post-optimization model (*Dopo*) across arrival rates ($\Delta D = 10\%$).

The breakdown confirms three qualitative shifts:

- **Mitigation of Premature Task Drops:** Under $BS = 2, BT = 2.0$ s, the optimized weight configuration substantially reduces the proportion of *Insufficient CPU* drops by forcing the solver to distribute tasks across available SENs rather than overloading localized nodes [19].
- **Elimination of Orbital Sunset Drops:** In the baseline runs (*Prima*), occasional task losses occurred due to satellites transitioning past the visibility horizon during routing. In the post-optimization runs (*Dopo*), drops categorized as *Sunset* are completely eliminated across all arrival rates.
- **Dominance of Physical Boundaries:** At high arrival rates ($\lambda \geq 8$ req/s), rejections in the optimized system are driven primarily by physical resource exhaustion (*Insufficient CPU+NET*) and strict deadline limits, verifying that the solver is operating at the physical capacity boundary of the constellation [6].

4.4.3 Response Time Containment

Figure 4.15 shows the response time decomposition comparison between the baseline and optimized configurations.




Figure 4.15. Response Time decomposition comparison: Pre-optimization baseline (*Prima*) vs. Post-optimization model (*Dopo*) showing containment of queuing latency ($\Delta D = 10\%$).

The latency decomposition highlights:

- **Containment of Queuing Latency:** For $BS = 2, BT = 2.0$ s, the total response time at peak load ($\lambda = 10$ req/s) is reduced from 1.12 ms in the baseline to 0.73 ms in the post-optimization model.
- **Invariance of Solver Runtime:** Across all test points, the pure solver computation time (T_{exec}) remains stable at ≈ 0.25 – 0.30 ms, confirming that increasing the primary weight to $\alpha = 1000$ does not increase the computational complexity of the HiGHS branch-and-cut routine.
- **Avoidance of Path Bloat:** By maintaining $\gamma = 8$, the post-optimization orchestrator avoids the path inflation observed with $\gamma \geq 1000$, ensuring that packets follow direct ISL trajectories without unnecessary multi-hop forwarding.

4.5 Summary of Findings and Architectural Readiness

The empirical evaluation and tuning campaign yield four foundational conclusions regarding centralized ILP orchestration in LEO satellite networks [6]:

1. **Queuing Dwell Time is the Dominant Bottleneck:** Solver computation time remains negligible ($T_{exec} \leq 0.60$ ms) across practical batch sizes. System degradation under heavy load is driven primarily by buffer residence times during batch aggregation, which trigger deadline violations before optimization begins.

2. **Necessity of Fast Dispatch:** To maintain service feasibility, centralized orchestrators must operate with aggressive timeouts ($BT \leq 0.05$ s) or minimal batch capacities ($BS \leq 2$), keeping buffer residence times well below application deadlines.
3. **Effectiveness of Objective Weight Calibration:** Enforcing a high primary performance weight ($\alpha = 1000$) focuses the solver on feasibility and deadline compliance, yielding up to +21% higher completion rates, while a moderate sunset penalty ($\gamma = 8$) eliminates connection drops due to orbital eclipse without inflating path lengths.
4. **Architectural Readiness for Multi-Algorithm Benchmarking:** The two calibrated configurations—reactive fast-dispatch ($BS = 20, BT = 0.05$ s) and controlled small-batching ($BS = 2, BT = 2.0$ s) with $\alpha = 1000$ and $\gamma = 8$ —represent the optimized operating envelope of the centralized ILP approach.

These fully tuned centralized configurations provide the baseline for Chapter 5, where their performance is benchmarked against state-of-the-art distributed heuristics (OrbitAware, DTS-base, DTS-APopt) and hierarchical ILP formulations across heterogeneous workload profiles [6].

Chapter 5

Comparative Benchmark

Having formalized the mathematical optimization framework in Chapter 3 and established the calibrated operational parameters in Chapter 4, this chapter presents an extensive comparative performance evaluation of the proposed Centralized Integer Linear Programming (ILP) Orchestrator against state-of-the-art distributed heuristics and hierarchical optimization baselines [6].

Deploying orbital edge intelligence across Low Earth Orbit (LEO) constellations involves balancing competing architectural paradigms [4, 1]. On the one hand, decentralized heuristics provide rapid, low-complexity decisions with negligible queuing overhead [6]. On the other hand, optimization models (both distributed and centralized) search the joint combinatorial space of compute allocation and multi-hop Inter-Satellite Link (ISL) routing to achieve theoretically optimal resource utilization [11, 6].

To quantify the operational trade-offs between these paradigms, this benchmark evaluates 9 algorithm configurations across a comprehensive testing matrix spanning diverse workload profiles, real-time payload dimensions, initial battery energy reserves, and request arrival rates.

5.1 Experimental Setup and Baseline Algorithms

The comparative evaluation is conducted within the event-driven discrete-event simulation framework developed in Python with `SimPy` [13], previously detailed in Section 4.1.2. The simulation faithfully reproduces satellite orbital kinematics [21], dynamic line-of-sight visibility windows, optical ISL transceivers [7, 8], processor queues, and onboard battery charge-discharge cycles [6].

5.1.1 Benchmark Dimensioning and Scenario Parameters

The orbital network is configured as a regional Walker-Delta LEO constellation operating at an orbital altitude of 600 km with an orbital inclination of 53° [16]. Adjacent satellites are interconnected via full-duplex optical ISLs characterized by nominal transmission bandwidths of $B_{u,v} = 1$ Gbps and dynamically computed propagation latencies $\tau_{u,v}$ [8]. Computational requests generated by Ground Users (GUs) are ingested through regional Ground Access Points (AP_{global}) and forwarded

to candidate Satellite Edge Nodes (SENs) [6].

The benchmark explores a multi-dimensional parameter space:

- **Workload Distribution Vector $\beta = (\beta_G, \beta_C, \beta_{CD})$:** Represents the proportion of Generic tasks, CPU-intensive tasks, and Data-intensive tasks [6]. We systematically evaluate three distinct operational regimes:
 1. *Mixed Multi-Service Profile*: $\beta = (0.40, 0.35, 0.25)$, representing a balanced heterogeneous mission [6].
 2. *Computation-Heavy Profile*: $\beta = (0, 1, 0)$, stressing onboard processor clock cycles [6].
 3. *Communication-Heavy Profile*: $\beta = (0, 0, 1)$, stressing ISL link capacity and data transfer energy [6].
- **Real-Time Payload Vector $\alpha = (\alpha_M, \alpha_L, \alpha_V)$:** Governs the distribution of input/output data payload sizes across Medium (1–5 MB), Large (5–20 MB), and Very Large (20–50 MB) files [6, 2]. Four payload profiles are benchmarked:
 1. *Standard (Std)*: $\alpha = (0.30, 0.50, 0.20)$ [6].
 2. *Large-heavy (L-heavy)*: $\alpha = (0, 0.70, 0.30)$.
 3. *Very-Large-heavy (VL-heavy)*: $\alpha = (0, 0.30, 0.70)$ [6].
 4. *Pure Large (Pure L)*: $\alpha = (0, 1.0, 0)$ [6].
- **Fleet Operational Energy Budget (E_i):** Candidate SENs are initialized with usable battery energy reserves of either 40 kJ (stringent energy constraint) or 80 kJ (relaxed energy constraint) per orbital pass [6, 17].
- **Traffic Intensity Spectrum (λ):** Request arrival rates vary systematically across $\lambda \in \{2, 3, 4, 6, 8, 10\}$ req/s, transitioning the network from a low-load regime (mean inter-arrival time of 0.50 s) to a heavy-load saturation regime (mean inter-arrival time of 0.10 s) [6].
- **Service Deadline Relaxation (ΔD):** In accordance with the macro-tuning findings of Chapter 4, the maximum allowable latency threshold incorporates a baseline relaxation factor of $\Delta D = 20\%$ to ensure fair evaluation under heavy traffic [6].

5.1.2 Evaluated Algorithms and State-of-the-Art Baselines

The comparative suite comprises 9 algorithmic configurations categorized into three architectural paradigms:

1. Distributed Online Heuristics:

- **DTS-base (Dynamic Task Scheduling Base):** A greedy distributed baseline that offloads tasks to the geographically closest available SEN possessing sufficient idle CPU capacity, without coordinating multi-hop ISL traffic [6].

- **DTS-Optimal (DTS Access-Point Optimized):** An enhanced multi-criteria greedy heuristic that jointly evaluates candidate SEN processing queues and the initial access-point ISL link bandwidth before dispatching tasks [6].
- **OrbitAware:** A trajectory-predictive distributed heuristic that leverages orbital ephemeris and contact window forecasting to select shortest routing paths and balance computational load dynamically [6].

2. Distributed / Hierarchical Optimization Models:

- **ILP-Hierarchical (Time):** A two-stage mathematical optimization formulation solved locally within satellite clusters, prioritizing task response time minimization [6].
- **ILP-Hierarchical (Energy):** A two-stage distributed formulation that prioritizes constellation energy preservation in the primary optimization phase [6].

3. Proposed Centralized ILP Orchestrator (Tuned Champions):

Following the micro-tuning calibration established in Chapter 4, the centralized orchestrator is evaluated under its optimal objective weights ($\alpha^* = 1000, \beta = 1.0, \gamma^* = 8$) across its two viable batching archetypes:

- **ILP-Centr (BS20, BT0.05, Time):** Reactive fast-dispatch configuration with a 50 ms timeout, optimizing for end-to-end response time.
- **ILP-Centr (BS20, BT0.05, Energy):** Reactive fast-dispatch configuration, optimizing for fleet cumulative energy preservation.
- **ILP-Centr (BS2, BT2, Time):** Controlled small-batching policy ($BS = 2$), optimizing for response time.
- **ILP-Centr (BS2, BT2, Energy):** Controlled small-batching policy ($BS = 2$), optimizing for fleet cumulative energy.

Table 5.1 provides a structured taxonomy of the evaluated scheduling algorithms.

Table 5.1. Taxonomy of evaluated task scheduling and routing algorithms in LEO satellite networks.

Algorithm	Paradigm	Decision Locus	Solver Complexity	State
DTS-base	Heuristic	Distributed / Local	$\mathcal{O}(S)$	Local
DTS-Optimal	Heuristic	Distributed / AP	$\mathcal{O}(S \log S)$	1-hop
OrbitAware	Heuristic	Distributed / Autonomous	$\mathcal{O}(S \log S)$	Ephem
ILP-Hierarchical (Time)	Optimization	Cluster Head	$\mathcal{O}(2^{ R_{sub} \cdot S_{sub} })$	Region
ILP-Hierarchical (Energy)	Optimization	Cluster Head	$\mathcal{O}(2^{ R_{sub} \cdot S_{sub} })$	Region
ILP-Centr (BS20, BT0.05, T/E)	Optimization	Centralized Master	$\mathcal{O}(2^{ R \cdot S })$ (Bounded)	Global
ILP-Centr (BS2, BT2, T/E)	Optimization	Centralized Master	$\mathcal{O}(2^{2 \cdot S })$	Global

5.2 Comparative Performance Evaluation across Workload Profiles

We evaluate the algorithms across three primary performance metrics:

1. **Task Success Rate (%)**: The fraction of total offloaded requests successfully routed, executed on an onboard SEN, and delivered to the requesting ground station within the deadline D_r without violating battery limits.
2. **End-to-End Response Time (ms)**: Decomposed into *System Execution Time* (T_{exec} , capturing solver resolution via HiGHS [20] and local CPU processing) and *Network/Wait Time* ($T_{net+wait}$, capturing batch queuing, ISL propagation, and transmission).
3. **Rejection Causes Breakdown (%)**: The diagnostic distribution of task drops categorized into *Deadline Exceeded*, *Insufficient CPU+NET*, *Insufficient CPU*, *No Server Found*, *ILP Infeasible*, and *Sunset*.

5.2.1 Mixed Multi-Service Workload Profile ($\beta = (0.40, 0.35, 0.25)$)

The mixed workload profile represents a standard multi-tenant operational scenario containing a heterogeneous combination of generic, CPU-heavy, and data-heavy requests [6]. Figures 5.1, 5.2, and 5.3 illustrate the Success Rate, Response Time decomposition, and Rejection Causes distribution, respectively.




Figure 5.1. Success Rate across arrival rates under Mixed Workload $\beta = (0.40, 0.35, 0.25)$ for 40 kJ (top) and 80 kJ (bottom) across payload distributions α .

Success Rate Dynamics

Under low arrival rates ($\lambda = 2$ req/s), all distributed heuristics (DTS-base, DTS-Optimal, OrbitAware) achieve exceptional completion rates exceeding 94–96% [6]. The proposed centralized orchestrator with fast-dispatch (ILP-Centr [BS20, BT0.05]) demonstrates competitive throughput, reaching 74–78%, while the small-batching configuration (ILP-Centr [BS2, BT2]) achieves 80–84% under Energy Mapping.




Figure 5.2. Response Time decomposition under Mixed Workload $\beta = (0.40, 0.35, 0.25)$ comparing System Execution and Network/Wait latency across arrival rates.




Figure 5.3. Rejection Causes distribution under Mixed Workload $\beta = (0.40, 0.35, 0.25)$ across arrival rates, battery budgets, and payload distributions.

As traffic intensity escalates toward heavy-load stress ($\lambda = 10$ req/s), the distributed heuristics exhibit superior resilience, maintaining completion rates between 78% and 83% [6]. Conversely, optimization-based approaches experience throughput degradation:

- ILP-Hierarchical models drop to 38–42% under 40 kJ battery budgets and recover modestly to 43–46% under 80 kJ [6].
- The centralized fast-dispatch variants (BS20, BT0.05) sustain completion rates of 40–44%, outperforming hierarchical formulations.
- The small-batching policy (BS2, BT2) declines to 30–36%, driven by buffer queuing accumulation during traffic bursts.

Response Time Analysis

Figure 5.2 highlights a clear structural divergence in latency components. Distributed heuristics achieve remarkably flat response times across all arrival rates, bounded

within 0.68–0.78 ms [6]. Because heuristics make instantaneous forwarding decisions, their queuing delay is virtually zero.

In contrast, optimization models display latency expansion under heavy load:

- **ILP-Hierarchical (Time)** incurs significant network and waiting inflation, with average response times climbing from 0.80 ms at $\lambda = 2$ req/s to 1.15–1.25 ms at $\lambda = 10$ req/s [6].
- Centralized configurations with fast dispatch (**BS20**, **BT0.05**) effectively contain latency growth, maintaining mean response times below 0.72 ms.
- Importantly, the pure computational execution time (T_{exec}) for the centralized ILP solver remains invariant and strictly bounded (< 0.30 ms), verifying that the branch-and-cut routine in HiGHS does not suffer from algorithmic stalling [20]. Latency growth is dominated entirely by the ingress buffering term $T_{net+wait}$.

Failure Diagnostics

Figure 5.3 reveals the physical nature of task drops in the mixed regime. Rejections for distributed heuristics are split evenly between *Insufficient CPU* (55–65%, green) and *Insufficient CPU+NET* (35–45%, orange) [6]. Because heuristics do not buffer requests, drops occur strictly when physical hardware or link transceivers are saturated.

Conversely, in the ILP-based models, *Deadline Exceeded* (blue) emerges as a substantial failure cause, accounting for 40–65% of all drops. This confirms that mathematical optimization models drop tasks primarily because ingress waiting time erodes the slack time required to complete execution within D_r . Finally, orbital *Sunset* drops remain strictly at 0% across all configurations, validating the robustness of the calibrated penalty $\gamma^* = 8$.

5.2.2 Computation-Heavy Stress: CPU-Intensive Profile ($\beta = (0, 1, 0)$)

When the workload consists entirely of compute-intensive requests ($\beta = (0, 1, 0)$), onboard processor clock cycles become the primary system bottleneck [6]. Figures 5.4, 5.5, and 5.6 display the performance metrics for this stress regime.

Success Rate Dynamics

The CPU-intensive profile induces severe performance differentiation. The distributed heuristics (**DTS-base**, **DTS-Optimal**, **OrbitAware**) demonstrate remarkable resilience, dropping gradually from 88–90% at $\lambda = 2$ req/s to 72–74% at $\lambda = 10$ req/s [6].

In sharp contrast, all optimization models experience steep degradation once traffic exceeds $\lambda = 4$ req/s:

- **ILP-Hierarchical** models drop from 75% at low load to 26–30% at peak arrival rates [6].
- Centralized fast-dispatch (**BS20**, **BT0.05**) maintains superior performance over hierarchical ILP, finishing at 38–42% under heavy load.




Figure 5.4. Success Rate across arrival rates under CPU-Intensive Workload $\beta = (0, 1, 0)$ for 40 kJ (top) and 80 kJ (bottom) across payload distributions α .




Figure 5.5. Response Time decomposition under CPU-Intensive Workload $\beta = (0, 1, 0)$ showing the dominance of queuing latency under heavy computational load.




Figure 5.6. Rejection Causes distribution under CPU-Intensive Workload $\beta = (0, 1, 0)$, illustrating the sharp dichotomy between heuristic drops and optimization deadline timeouts.

- Centralized small-batching (BS2, BT2) degrades to 32–36% at $\lambda = 10$ req/s.
- Interestingly, increasing the battery energy budget from 40 kJ to 80 kJ produces negligible throughput improvements ($< 1.5\%$), proving that in CPU-bound regimes, battery capacity is not an active constraint [6].

The Fundamental Rejection Dichotomy

Figure 5.6 presents the most profound architectural insight of the benchmark. We observe a complete, binary divergence in failure causes between heuristics and optimization formulations:

- **Distributed Heuristics (DTS, OrbitAware):** Exactly 100% of all task rejections are caused by *Insufficient CPU* (green bars) [6]. Because heuristics evaluate nodes greedily, if a satellite’s CPU core is occupied, the task is dropped immediately without waiting. Consequently, not a single task fails due to *Deadline Exceeded*.
- **Mathematical Optimization Models (ILP-Hierarchical, ILP-Centralized):** Between 60% and 88% of all rejections are caused by *Deadline Exceeded* (blue bars). Under heavy computational demands, processor queues ($W_{cpu,i}$) expand rapidly across all candidate satellites. When the ILP solver aggregates tasks, the mathematical constraint:

$$R_{r,i} = T_{up}(r, i) + W_{cpu,i} + \frac{d_{cpu,r}}{C_{sen,i}} \leq D_r - T_{wait} \quad (5.1)$$

becomes infeasible for a majority of candidate nodes. Tasks fail not because the constellation lacks physical transceivers, but because cumulative queuing delay violates the rigid application deadline.

5.2.3 Communication-Heavy Stress: Data-Intensive Profile ($\beta = (0, 0, 1)$)

In the data-intensive regime ($\beta = (0, 0, 1)$), tasks carry minimal CPU requirements but demand the transfer of large payload files (20–50 MB) across multi-hop ISL routes [6]. Figures 5.7, 5.8, and 5.9 display the comparative performance metrics.

Success Rate and Global Optimality Supremacy

The data-intensive scenario produces a remarkable inversion in algorithmic ranking:

- At low to moderate arrival rates ($\lambda \leq 3$ req/s), mathematical optimization models decisively outperform distributed heuristics. Under 80 kJ battery reserves, ILP-Hierarchical (Time) achieves a near-perfect Success Rate of 98–100%, surpassing DTS-base (93%) and OrbitAware (94%) [6].
- The centralized orchestrator (ILP-Centr [BS20, BT0.05]) achieves robust throughput, maintaining 70–74% completion at $\lambda = 2$ req/s and outperforming small-batching.




Figure 5.7. Success Rate across arrival rates under Data-Intensive Workload $\beta = (0, 0, 1)$ for 40 kJ (top) and 80 kJ (bottom) across payload distributions α .




Figure 5.8. Response Time decomposition under Data-Intensive Workload $\beta = (0, 0, 1)$ across arrival rates and battery capacities.




Figure 5.9. Rejection Causes distribution under Data-Intensive Workload $\beta = (0, 0, 1)$, showing near-total dominance of physical ISL and energy saturation (*Insuff. CPU+NET*).

- **Critical Impact of Energy Budget (40 kJ vs. 80 kJ):** Unlike the CPU-intensive case, the data-intensive regime is highly sensitive to the constellation battery capacity [6]. Moving large payloads across optical links incurs heavy transceiver power consumption ($P_{net} \cdot S_r / B_{u,v}$). At 40 kJ, satellites exhaust their orbital energy budget early, causing success rates to drop to 48–52% at peak traffic. Doubling the battery budget to 80 kJ restores completion rates to 82–85% for ILP-Hierarchical and improves centralized performance by +6–9%.

Structural Uniformity of Rejection Causes

Figure 5.9 displays striking uniformity across all 9 algorithms: **over 96–99% of all task rejections are categorized as *Insufficient CPU+NET*** (orange bars). In this regime, CPU cores remain idle while optical ISL channel bandwidth and node transmission power are pushed to physical saturation. Because the bottleneck is purely infrastructural (link transceiver capacity and battery limits), global mathematical routing paths provide tangible benefits by steering payloads around saturated inter-plane links.

5.3 Trade-Off Analysis and Operational Discussion

Correlating the empirical findings across the 9 evaluated scenarios synthesizes four foundational operational trade-offs governing satellite task offloading [6].

5.3.1 Theoretical Global Optimality vs. Ingress Queuing Latency

The central dilemma of centralized satellite orchestration lies in the trade-off between global scheduling efficiency and ingress queuing latency [6]:

- **Distributed Heuristics:** Operate on an $M/M/1$ or $M/G/1$ queue with zero buffer aggregation. Tasks are evaluated immediately upon arrival at the gateway Access Point. This responsiveness yields exceptional resilience against heavy request arrival rates ($\lambda > 6$ req/s), preventing deadline expiration [6]. However, lacking global visibility, heuristics make greedy, localized decisions that cause uneven battery drain and cannot avoid downstream link congestion.
- **Centralized ILP:** Solves an exact combinatorial problem over the entire constellation graph [11]. While the solver runtime is negligible ($T_{exec} \leq 0.30$ ms), the necessity of buffering requests to form batches introduces static ingress delay (T_{wait}). In CPU-intensive scenarios, this waiting time consumes valuable deadline slack, triggering cascade rejections.

5.3.2 Energy Budget Sensitivity Across Computational Dimensions

Our results demonstrate that battery reserve sizing (40 kJ vs. 80 kJ) exerts an asymmetric impact depending on workload dimensionality:

- **Compute-Bound Regimes ($\beta = (0, 1, 0)$):** Inelastic to energy scaling. The system bottleneck is governed strictly by the CPU clock frequency ($C_{sen,i} =$

1 GHz) and maximum concurrent task limits [6]. Adding battery reserves does not prevent processor queue overflow.

- **Data-Bound Regimes ($\beta = (0, 0, 1)$):** Highly elastic to energy scaling. Moving 50 MB payloads across 4 ISL hops consumes several kilojoules of transmission energy. Expanding battery capacity from 40 kJ to 80 kJ unlocks substantial throughput gains (+15–25%), enabling sustained multi-hop routing.

5.3.3 Fast-Dispatch ($BT = 0.05$ s) vs. Controlled Small-Batching ($BS = 2$)

Comparing the two centralized archetypes validates the critical importance of dispatch responsiveness:

- The **Fast-Dispatch configuration** ($BS = 20, BT = 0.05$ s) consistently outperforms small-batching under heavy load ($\lambda \geq 6$ req/s). By enforcing an aggressive 50 ms timeout, the coordinator flushes tasks almost instantaneously, effectively emulating an online pipeline while retaining mathematical constraint checking.
- The **Controlled Small-Batching configuration** ($BS = 2, BT = 2.0$ s) excels at low traffic ($\lambda \leq 3$ req/s), where waiting for a second task provides the solver with multi-request optimization opportunities without risking deadline expiration.

5.3.4 Architectural Synthesis and Deployment Guidelines

Based on the empirical evidence, Table 5.2 outlines prescriptive architectural guidelines recommending the optimal scheduling strategy based on mission profile and traffic conditions.

Table 5.2. Prescriptive architectural deployment matrix for Orbital Edge Computing missions.

Workload Characteristics	Traffic Regime	Recommended Paradigm	Primary Rationale
Data-Intensive (Earth Observation)	Low to Moderate	ILP-Hierarchical / Centr	Maximizes ISL route
Data-Intensive (Large Payloads)	Heavy Stress	DTS-Optimal + Large Battery	Avoids ingress queui
CPU-Intensive (AI Inference)	Any Traffic	OrbitAware / DTS-base	Zero queuing; immed
Mixed Heterogeneous Services	$\lambda \leq 4$ req/s	ILP-Centr (BS2, BT2)	Balances joint alloca
Mixed Heterogeneous Services	$\lambda > 6$ req/s	ILP-Centr (BS20, BT0.05)	Low buffer dwell tim
Critical Emergency Alerts	Ultra-Low Latency	Distributed Heuristic	Minimizes single-poi

5.4 Summary of Chapter Findings

The comprehensive comparative benchmark establishes four definitive conclusions:

1. **Validation of Fast-Dispatch Centralized Scheduling:** The tuned centralized orchestrator operating under fast-dispatch ($BS = 20, BT = 0.05$ s)

maintains competitive completion rates (74–78% at low traffic, 40–44% under stress) while strictly bounding solver computation latency ($T_{exec} < 0.30$ ms).

2. **Isolation of the CPU Deadline Bottleneck:** In compute-heavy regimes, optimization-based failures are overwhelmingly driven by temporal deadline expiration (60–88% *Deadline Exceeded*), whereas heuristic failures stem strictly from physical hardware limits (100% *Insufficient CPU*).
3. **Energy Elasticity in Communication Regimes:** Data-intensive workloads reward global mathematical routing, enabling ILP formulations to achieve up to 100% success rates when supported by sufficient battery reserves (80 kJ).
4. **Complete Elimination of Sunset Loss:** Across all 9 evaluated configurations and hundreds of simulation runs, connection drops due to orbital eclipse remained at exactly 0%, confirming the absolute efficacy of the calibrated multi-objective sunset penalty ($\gamma^* = 8$).

These empirical findings define the operational boundaries of centralized orbital orchestration, providing the foundation for the concluding assessments and future research directions presented in Chapter 6.

Chapter 6

Conclusions and Future Work

This thesis investigated the operational feasibility, performance boundaries, and systemic trade-offs of centralized mathematical optimization for computational task offloading and multi-hop routing in Low Earth Orbit (LEO) satellite networks [1, 6]. By designing, formalizing, and implementing a Centralized Integer Linear Programming (ILP) Orchestrator coupled with an event-driven dynamic batching mechanism within a `SimPy` discrete-event simulation environment [13], this work systematically explored the viability of orbital edge intelligence under realistic spaceborne constraints [4, 6].

This concluding chapter synthesizes the primary findings and answers the core research question formulated at the outset of the thesis. Furthermore, it details the critical physical and architectural limitations inherent to centralized space computing and outlines promising future research directions toward resilient, hybrid orbital orchestration paradigms [5, 1].

6.1 Synthesis of Contributions and Research Question Resolution

6.1.1 Answering the Core Research Question

In Section 1.3, this thesis posed the foundational research question: *To what extent can a centralized optimization approach maintain global allocation efficiency in high-density satellite networks before queuing and solver overhead cause systemic performance degradation?*

Based on the multi-tiered experimental evaluations conducted across Chapters 4 and 5, the answer is clear: **the performance ceiling of centralized orbital orchestration is fundamentally architectural rather than computational.**

- **Computational Tractability of the Solver:** Contrary to the conventional assumption that NP-hard mathematical optimization is prohibitively slow for real-time spaceborne execution, the open-source `HiGHS` branch-and-cut solver demonstrated remarkable computational efficiency [20]. By bounding the active optimization batch size ($BS \leq 20$), the decision variable space remained strictly constrained, maintaining pure solver execution times at $T_{exec} \leq 0.30\text{--}0.60$ ms across all tested arrival rates ($\lambda \in [2, 10]$ req/s). This

proves that solving regional linear formulations is computationally viable on contemporary space-qualified onboard processing units [5].

- **Ingress Queuing as the Systemic Bottleneck:** System throughput collapse under elevated arrival rates ($\lambda \geq 6$ req/s) is governed almost entirely by the ingress queuing latency (T_{wait}) accumulated within the orchestrator buffer prior to solver invocation. As traffic scales, request aggregation delays severely erode task slack margins ($\Delta t_{rem} = D_r - t_{queue}$), triggering cascade *Deadline Exceeded* dropouts before the scheduling algorithm can be executed.
- **Primacy of Fast-Dispatch Policies:** To prevent queue-induced deadline collapse, centralized orchestrators must adopt responsive dispatch mechanisms. Operating with an aggressive timeout ($BS = 20, BT \leq 0.05$ s) maintains near-zero buffer residence times, effectively transforming the batching pipeline into a semi-continuous dispatch engine that sustains competitive completion rates (74%–78% at low traffic, 40%–44% under stress).

6.1.2 Algorithmic Archetypes and Workload Sensitivity

The comparative benchmark in Chapter 5 established distinct operational envelopes across heterogeneous application profiles [6]:

- **Computation-Heavy Workloads ($\beta = (0, 1, 0)$):** Distributed online heuristics (*OrbitAware*, *DTS-base*, *DTS-Optimal*) demonstrate superior resilience, achieving 72%–74% success rates at peak load ($\lambda = 10$ req/s) [6]. Crucially, the rejection diagnostics revealed a complete operational divide: heuristics dropped 100% of rejected tasks due to physical capacity exhaustion (*Insufficient CPU*), whereas ILP models suffered 60%–88% of their failures due to *Deadline Exceeded*. In compute-bound environments, instantaneous local assignment without ingress buffering is essential to prevent temporal expiration.
- **Communication-Heavy Workloads ($\beta = (0, 0, 1)$):** The centralized optimization paradigm achieves clear superiority in data-intensive environments. When offloading large datasets (20–50 MB), global visibility of multi-hop Inter-Satellite Link (ISL) states enabled the ILP formulations to attain near-perfect completion rates (98%–100%) under low-to-moderate loads, outperforming greedy heuristics. Furthermore, data-intensive workloads exhibited strong elasticity with respect to satellite battery capacity: expanding battery reserves from 40 kJ to 80 kJ yielded a +15%–25% throughput recovery, proving that mathematically coordinated routing paths preserve energy across optical transceivers [8, 6].
- **Multi-Objective Weight Calibration ($\alpha^* = 1000, \gamma^* = 8$):** The micro-tuning campaign proved that dominating the primary objective ($\alpha^* = 1000$) ensures that solver decisions are driven strictly by deadline compliance and capacity limits, yielding an average +21% gain in success rate. Simultaneously, calibrating the orbital eclipse penalty to $\gamma^* = 8$ completely eliminated *Sunset* drops (0% loss across all runs) while avoiding the 55% latency inflation and

topological path-bloat that occurs when the eclipse penalty is over-weighted ($\gamma \geq 1000$).

6.2 Critical Limitations of the Centralized Approach

Despite the mathematical rigor and global optimality demonstrated by the proposed framework, deploying a pure centralized orchestrator in operational mega-constellations involves severe structural limitations [5, 1].

6.2.1 Ingress Buffering and the Batching Dilemma

Centralized mathematical scheduling requires aggregating incoming discrete requests into batches to instantiate the linear program [11]. This creates an inescapable trade-off between optimization horizon and buffering latency:

$$T_{total}(r) = T_{wait}(BS, BT, \lambda) + T_{up}(r, i) + W_{cpu,i} + \frac{d_{cpu,r}}{C_{sen,i}} + T_{down}(i, AP_r) \quad (6.1)$$

While larger batch sizes ($BS \geq 8$) allow the branch-and-cut algorithm to uncover superior combinatorial placement patterns across the fleet, the static wait time T_{wait} directly consumes the application deadline slack. For mission-critical real-time services characterized by tight latency bounds ($D_r \leq 1.0$ s), any artificial queuing introduced before solver execution severely degrades service availability [2].

6.2.2 Single Point of Failure and Space Hazards

Relying on a designated orbital Master node introduces an acute Single Point of Failure (SPOF) [5]. Spaceborne edge computing operates in an exceptionally harsh environment subject to extreme thermal cycling, total ionizing dose (TID), and single-event upsets (SEUs) induced by cosmic rays and solar particle events [4, 5]. If the central Master satellite suffers a radiation-induced latch-up, an onboard computer reboot, an attitude control anomaly, or requires an autonomous orbital collision avoidance maneuver, the entire task scheduling and forwarding pipeline across the constellation is immediately paralyzed.

6.2.3 State Telemetry Overhead and Information Staleness

To formulate an accurate mathematical representation of constellation constraints, the orchestrator must maintain an updated global state matrix containing processor queue lengths ($W_{cpu,i}$), residual battery levels (B_i), and active ISL bandwidths ($B_{u,v}$). Transmitting this telemetry across dynamic inter-satellite links incurs control signaling overhead that scales with the number of satellites ($\mathcal{O}(|S|)$) [1].

Due to the extreme orbital velocities of LEO spacecraft (~ 7.5 km/s), continuous relative motion alters line-of-sight angles and link propagation delays [7]. Consequently, by the time state telemetry is gathered, buffered, and solved, the physical state of the constellation may have drifted. This state staleness risks rendering calculated schedules sub-optimal or infeasible mid-execution, necessitating reliance on emergency greedy fallback routines.

6.3 Future Research Directions

The conclusions and limitations identified in this thesis pave the way for several natural extensions and advanced research trajectories.

6.3.1 Two-Tier Hybrid Orchestration Architectures

To overcome the trade-off between centralized optimality and distributed responsiveness, future architectures should investigate two-tier hierarchical frameworks:

- **Tier 1 (Fast-Path Heuristic Layer):** Time-critical and CPU-intensive requests are evaluated immediately at the local Access Point using decentralized, zero-queuing heuristics (such as `OrbitAware` or `DTS-Optimal`), ensuring sub-millisecond dispatching [6].
- **Tier 2 (Global Optimization Layer):** Computationally relaxed, data-intensive tasks (such as delay-tolerant Earth Observation image filtering) are buffered and offloaded to localized cluster-head ILP solvers [9, 15]. This dual-track paradigm harnesses mathematical efficiency for bandwidth-heavy traffic while preserving ultra-low latency for time-sensitive services.

6.3.2 Traffic-Aware Adaptive Batching via Reinforcement Learning

The static batching thresholds (BS, BT) evaluated in this work could be replaced by an intelligent, dynamic batching engine. A lightweight Reinforcement Learning (RL) agent—such as a contextual bandit or Q-learning controller—could be deployed at the ingress gateway to dynamically modulate the timeout threshold $BT(t)$ in real time [10]:

$$BT(t) = f\left(\hat{\lambda}(t), \min_{r \in Q}\{D_r - t_{arr,r}\}, \bar{W}_{cpu}(t)\right) \quad (6.2)$$

By tracking traffic arrival surges ($\hat{\lambda}$) and the minimum residual slack time of buffered requests, the orchestrator could trigger immediate flushes during high-risk intervals and expand batch sizes during traffic lulls, optimizing the trade-off between scheduling efficiency and deadline preservation.

6.3.3 Multi-Agent Distributed Reinforcement Learning (MADRL)

To eliminate the centralized Single Point of Failure and bypass the communication overhead of global telemetry aggregation, decentralized Multi-Agent Deep Reinforcement Learning (MADRL) presents a compelling alternative [10]. Equipping each Satellite Edge Node with an autonomous actor-critic policy trained under Centralized Training with Decentralized Execution (CTDE) would allow satellites to make collaborative routing and scheduling decisions based solely on localized, one-hop neighborhood observations [10]. This model could approximate the global efficiency of ILP solutions while achieving the instantaneous decision speeds characteristic of distributed heuristics.

6.3.4 Hardware-in-the-Loop (HIL) Prototyping on Space-Qualified Hardware

Finally, extending this simulation-based validation toward physical implementation requires testing the optimization pipeline on physical embedded hardware architectures. Deploying the HiGHS solver [20] and the batching engine on radiation-hardened system-on-chip (SoC) architectures—such as the ARM Cortex-R52, Cortex-A53, or fault-tolerant space-grade RISC-V platforms (e.g., NOEL-V)—will provide empirical measurements of CPU instruction overhead, memory footprint, cache contention, and power consumption during real-time solver execution [5].

6.4 Concluding Remarks

The transition from passive bent-pipe relaying to active Orbital Edge Computing marks a significant evolution in space-based information systems [4, 1]. This thesis demonstrated that centralized Integer Linear Programming is computationally viable on spaceborne processing platforms, effectively bounding solver execution latency within sub-millisecond limits [20]. However, the temporal latency imposed by request buffering in ingress queues represents a fundamental physical boundary that limits the scalability of centralized scheduling under high-frequency traffic.

By quantifying the operational trade-offs between centralized mathematical optimization, distributed hierarchical solvers, and autonomous heuristics across diverse mission profiles, this work provides architectural guidelines for designing the next generation of resilient, energy-aware orbital computing systems [6].

Bibliography

- [1] Y. Zhang, X. Li, H. Zhou, and D. Chen, “A comprehensive survey on orbital edge computing: Systems, applications, and algorithms,” *IEEE Communications Surveys & Tutorials*, vol. 25, no. 4, pp. 2100–2135, 2023.
- [2] I. Leyva-Mayorga, M. Martinez-Gost, M. Moretti, A. Pérez-Neira, M. Vázquez, P. Popovski, and B. Soret, “Satellite edge computing for real-time and very-high resolution earth observation,” *IEEE Transactions on Communications*, vol. 71, no. 10, pp. 6180–6194, 2023.
- [3] S. Wang, Q. Zhang, C. Zhao *et al.*, “Satellite computing: Vision and challenges – Tiansuan constellation,” *IEEE Internet of Things Journal*, vol. 10, no. 20, pp. 17 550–17 563, 2023.
- [4] B. Denby and B. Lucia, “Orbital edge computing: Nanosatellite clusters as accelerators for earth observation,” *IEEE Micro*, vol. 40, no. 1, pp. 13–21, 2019.
- [5] —, “Orbital edge computing: Nanosatellite constellations as a new class of computer system,” in *Proceedings of the 25th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, 2020, pp. 939–954.
- [6] E. Casalicchio *et al.*, “Handling of energy- and time-constrained heterogeneous satellite computing services,” in *Proceedings of the IEEE International Conference on Distributed Computing Systems (ICDCS)*. IEEE, 2025, pp. 1–12.
- [7] Q. Chen, G. Giambene, L. Yang, C. Fan, and X. Chen, “Analysis of inter-satellite link paths for LEO mega-constellation networks,” *IEEE Transactions on Vehicular Technology*, vol. 70, no. 3, pp. 2743–2755, 2021.
- [8] A. U. Chaudhry and H. Yanikomeroglu, “Laser intersatellite links in a Star-link constellation: A classification and analysis,” *IEEE Vehicular Technology Magazine*, vol. 16, no. 2, pp. 48–56, 2021.
- [9] L. Zhong, Y. Li, M.-F. Ge, M. Feng, and S. Mao, “Joint task offloading and resource allocation for LEO satellite-based mobile edge computing systems with heterogeneous task demands,” *IEEE Transactions on Vehicular Technology*, vol. 74, no. 7, pp. 11 337–11 352, 2025.
- [10] W. Fan, Q. Meng, G. Wang, H. Bian, and Y. Liu, “Satellite edge intelligence: DRL-based resource management for task inference in LEO-based satellite-

- ground collaborative networks,” *IEEE Transactions on Mobile Computing*, vol. 24, no. 10, pp. 10 710–10 728, 2025.
- [11] L. A. Wolsey, *Integer Programming*. New York, NY, USA: John Wiley & Sons, 1998.
- [12] S. Boyd and L. Vandenberghe, *Convex Optimization*. Cambridge, UK: Cambridge University Press, 2004.
- [13] S. Scherfke and O. Lünsdorf, *SimPy: Discrete Event Simulation for Python*, 2023, version 4.1.1. [Online]. Available: <https://simpy.readthedocs.io/>
- [14] S. Mitchell, M. O’Sullivan, and A. Dambrosio, “Pulp: A linear programming toolkit for Python,” The University of Auckland, Tech. Rep., 2011, version 2.8.
- [15] Q. Jiang, L. Zheng, Y. Zhou, H. Liu, Q. Kong, Y. Zhang, and B. Chen, “Efficient on-orbit remote sensing imagery processing via satellite edge computing resource scheduling optimization,” *IEEE Transactions on Geoscience and Remote Sensing*, vol. 63, pp. 1–19, 2025.
- [16] J. G. Walker, “Satellite constellations,” *Journal of the British Interplanetary Society*, vol. 37, pp. 559–571, 1984.
- [17] P. Peng, T. Xu, X. Chen, C. C. Zarakovitis, and C. Wu, “Blocked-job-offloading-based computing resources sharing in LEO satellite networks,” *IEEE Internet of Things Journal*, vol. 12, no. 2, pp. 2287–2290, 2025.
- [18] A. Neumann, C. Aichele, M. Lindner, and S. Fux, “Better approach to mobile ad-hoc networking (B.A.T.M.A.N.),” IETF Draft, Tech. Rep., 2008, draft-wunderlich-openmesh-manet-routing-00.
- [19] X. Yang, H. Dai, Y. Tang, and X. Yi, “A general metric of load balancing in δ -range,” in *Proceedings of Advanced Parallel Processing Technologies (APPT)*. Springer, 2003, pp. 311–321.
- [20] Q. Huangfu and J. A. J. Hall, “Parallelizing the dual revised simplex method,” *Mathematical Programming Computation*, vol. 10, no. 1, pp. 119–142, 2018.
- [21] D. Vallado, P. Crawford, R. Hujsak, and T. S. Kelso, “Revisiting spacetrack report no. 3: Rev 2,” in *AIAA/AAS Astrodynamics Specialist Conference and Exhibit*, 2006, p. 6753.
- [22] E. W. Dijkstra, “A note on two problems in connexion with graphs,” *Numerische Mathematik*, vol. 1, no. 1, pp. 269–271, 1959.
- [23] C. R. Harris, K. J. Millman, S. J. van der Walt, R. Gommers, P. Virtanen, D. Cournapeau, E. Wieser, J. Taylor, S. Berg, N. J. Smith *et al.*, “Array programming with NumPy,” *Nature*, vol. 585, no. 7825, pp. 357–362, 2020.
- [24] P. Virtanen, R. Gommers, T. E. Oliphant, M. Haberland, T. Reddy, D. Cournapeau, E. Burovski, P. Peterson, W. Weckesser, J. Bright *et al.*, “SciPy 1.0: Fundamental algorithms for scientific computing in Python,” *Nature Methods*, vol. 17, no. 3, pp. 261–272, 2020.

-
- [25] J. D. Hunter, “Matplotlib: A 2d graphics environment,” *Computing in Science & Engineering*, vol. 9, no. 3, pp. 90–95, 2007.